

# Why Graphical Models and Bayesian methods? – 1

- formalization of *Information Processing*
  - data is information
  - sensors give information
  - outputs/actions/decisions are *missing* information (to be ‘inferred’)
  - coupling between sources/points of information

⇒ Graphical Models formalize “networks of coupled information”

⇒ Information Processing can be viewed as inference or message passing in Graphical Models

# probability theory

- why do we need probabilities?
  - of course, in case of random events, stochasticity...
- but also in a deterministic world!:
  - lack of knowledge!
  - hidden (latent) variables
  - expressing *uncertainty*
  - expressing *information*
- probabilities are a generic tool to express uncertainty, information, and coupling

# Probability: Frequentist and Bayesian

- Frequentist probabilities are defined in the limit of an infinite number of trials
- *Example:* The probability of a particular coin landing heads up is 0.43
- Bayesian (subjective) probabilities quantify degrees of belief
- *Example:* The probability of it raining tomorrow is 0.3
- Not possible to repeat tomorrow many times

# random variables

- intuitively: a *random variable* takes on *values* with a certain probability  
a bit more formally: a random variable relates a measurable space with a domain (sample space) and thereby introduces a probability measure on the domain (“assigns a probability to each possible value”)
- the *domain*  $\text{dom}(X)$  of a variable  $X$  is the set possible values of a random variable (mutually exclusive and collectively exhaustive)  
*Example:* a dice can take values  $\{1, \dots, 6\}$
- we use capital letters  $X$  to denote random variables and lower case letters  $x$  to denote values that they take
- we use the  $P$  to denote the mapping to probabilities

# random variables (in terms of sets)

Let  $X$  be a random variable with domain  $\Omega = \text{dom}(X)$

Let  $A, B \subset \Omega$  be subsets of the domain and  $x \in \Omega$  a value in the domain.

- $X \in A$  or  $X \in B$  or  $X = x$  are called *events*
  - we use the  $P$  to denote the mapping to probabilities:
    - $P(X \in A) \in \mathbb{R}$
  - we require
    - $P(X \in \emptyset) = 0$  and  $P(X \in \Omega) = 1$
    - if  $A \cap B = \emptyset$  then  $P(X \in A \cup B) = P(X \in A) + P(X \in B)$
- if the domain is discrete this implies *normalization*:
- $\sum_{x \in \Omega} P(X = x) = 1$

# probability distribution & tables

- for continuous domains: “probability distribution” is the integral of a “probability density function”
- for discrete domains: “probability distribution” and “probability mass function” are used synonymously

- a RV assigns a probability to each possible value  
→ think of the probability distribution as a *table* of numbers:

*Example:* A fair dice  $X$ ,  $\text{dom}(X) = \{1, 2, 3, 4, 5, 6\}$ , with

$$\forall_{x \in \text{dom}(X)} : P(X = x) = \frac{1}{6}$$

corresponds to the table

$$\left[ \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6} \right]$$

- in implementations we typically represent random variables by tables (arrays/vectors) of numbers

# joint distributions

- assume we have two random variable  $X$  and  $Y$ . The *joint probability distribution*

$$P(X = x, Y = y)$$

gives the probability that  $X = x$  *and*  $Y = y$ .

(In logic one would perhaps write something like  $X = x \wedge Y = y$ . But not so in joint probability distributions.)

- *Example:* Suppose Toothache and Cavity are the variables:

	Toothache = true	Toothache = false
Cavity = true	0.04	0.06
Cavity = false	0.01	0.89

we write

$$P(\text{Toothache} = \text{true}, \text{Cavity} = \text{false}) = 0.01$$

# joint distributions

- note, most of what'll need will be about JOINT PROBABILITY DISTRIBUTIONS
  - graphical models are nothing but descriptions of joint probability distributions!
  - correlations, interdependence, coupling are all expressed in terms of joint probability distributions
  - whenever you're confused about the “model”, the “approach”, the “assumptions”, etc, reconsider explicitly what the joint probability distribution over all involved variables is!



# joint distributions

- *definitions:*

- the *marginal* (probability) of  $X$  given  $P(X, Y)$  is

$$P(X) = \sum_Y P(X, Y)$$

- the *conditional* (probability) of  $X$  given  $Y$  and  $P(X, Y)$  is

$$P(X|Y) = \frac{P(X, Y)}{P(Y)}$$

defs also hold for tuples of variables, e.g.,  $X = (X_1, \dots, X_n)$ ,  $Y = (Y_1, \dots, Y_m)$

- *implications:*

- the *product rule*  $P(X, Y) = P(X|Y) P(Y) = P(Y|X) P(X)$

- the *chain rule*  $P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i | X_1, \dots, X_{i-1})$

- *Bayes Rule*

$$P(X|Y) = \frac{P(Y|X)}{P(Y)} P(X)$$

Ich gehe nach Hause....  
 $P(\text{Ich}) P(\text{gehe} | \text{Ich}) P(\text{nach} | \text{Ich})$

# Bayes Rule

$P(\text{Hause} | \text{Hause})$   ~~$P(\text{Hause} | \text{Hause})$~~

- Thomas Bayes (1702–1761)
- Bayes Rule is a trivial implication of the definitions of marginal and conditional probability!

$$P(Y|X) = \frac{P(X,Y)}{P(X)}$$

- importance lies in its interpretation and use:

$$P(X|Y) = \frac{P(Y|X)}{P(Y)} P(X), \quad \text{posterior} = \frac{\text{likelihood}}{\text{evidence}} \text{prior}$$

$$P(\text{cause}|\text{effect}) = \frac{P(\text{effect}|\text{cause})}{P(\text{effect})} P(\text{cause})$$

*Example:* let  $M$  be meningitis,  $S$  be stiff neck

$$P(M|S) = \frac{P(S|M)}{P(S)} P(M) = \frac{0.8}{0.1} 0.0001 = 0.0008$$

Note: posterior probability of meningitis still very small

N-Gram: (shingles)  
↑

Ich gehe nach Hause

1-Gram / Unigram: "ich", "gehe", "nach" ...  
1 2 3

2-Gram: "ich gehe", "gehe nach", "nach Hause"  
1 2 3

3-Gram: "ich gehe nach", "gehe nach Hause"

# inference

- we will deal with many variables

$$X = (H_1, \dots, H_n, E_1, \dots, E_m, Y_1, \dots, Y_k)$$

- we are given the joint probability distribution

$$P(H_1, \dots, H_n, E_1, \dots, E_m, Y_1, \dots, Y_k)$$

- some variables  $E_1, \dots, E_m$  are observed (we have evidence)  
for the other variables  $H_1, \dots, H_n, Y_1, \dots, Y_k$  we have no evidence  
we want to know the *posterior* over some variables  $Y_1, \dots, Y_K$

$$P(Y_{1:k} \mid E_{1:m}) = \frac{P(Y_{1:k}, E_{1:m})}{P(E_{1:m})} \propto \sum_{H_{1:n}} P(Y_{1:k}, E_{1:m}, H_{1:n}) \quad (1)$$

- computing  $P(Y_{1:k} \mid E_{1:m})$  is the *problem of inference*
- obvious problem: size of table  $P(Y_{1:k}, E_{1:m}, H_{1:n})$  is  $d^{k+m+n}$

# summary

- focus of this lecture:
  - graphical models as a generic tool for inference with coupled random variables
  - probability theory as calculus for uncertainty, information, evidence
  - learning graphical models from data
  - using graphical models for decision making & RL
- next time:
  - naive Bayes
  - graphical models
  - inference using the elimination algorithm

# cheat sheat

- a random variable  $X$  assigns probabilities  $P(X=x) \in \mathbb{R}$  to values  $x \in \text{dom}(x)$
- probability distribution  $\leftrightarrow$  table (vector) of probabilities for each value  
(normalization:  $\sum_X P(X) = 1$ )
- joint distribution  $P(X, Y) \leftrightarrow$  table (matrix) of probabilities
- definition: marginal  $P(X) = \sum_Y P(X, Y)$  (summing along columns/rows)
- definition: conditional  $P(X|Y) = \frac{P(X,Y)}{P(Y)}$  (normalizing each column)
- implications:

$$P(X, Y) = P(X|Y) P(Y) = P(Y|X) P(X)$$

$$P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i | X_1, \dots, X_{i-1})$$

$$P(X|Y) = \frac{P(Y|X)}{P(Y)} P(X), \quad \text{posterior} = \frac{\text{likelihood}}{\text{evidence}} \text{prior}$$

- definition: *inference* is the problem to compute

$$P(Y_{1:k} \mid E_{1:m}) = \frac{P(Y_{1:k}, E_{1:m})}{P(E_{1:m})} \propto \sum_{H_{1:n}} P(Y_{1:k}, E_{1:m}, H_{1:n})$$

- **web links:**

Bayes Rule:

<http://www.cs.ubc.ca/~murphyk/Bayes/bayesrule.html>

Kevin's lecture:

[http://www.cs.ubc.ca/~murphyk/Teaching/CS532c\\_Fall10/Lectures/index.html](http://www.cs.ubc.ca/~murphyk/Teaching/CS532c_Fall10/Lectures/index.html)

<http://www.cs.ubc.ca/~murphyk/Bayes/bnsoft.html>

site: <http://www.cs.ubc.ca/~murphyk/Bayes>

# Overview

- graphical models
  - Bayesian networks
  - Markov random fields
- ~~inference~~
  - ~~belief propagation~~
  - ~~loopy belief propagation~~
- assumption:
  - graph structure is known
  - probability tables are known
  - realistic?

Klaus	chase	Karl-Heinz
N	✓	N
P	∅	P



# Learning

- Nomenclature
  - Input variables / observations:  $x$
  - Output variables / targets:  $y$
- Recall:  $P(y|X = x) = P(x|y)P(y)/P(x)$
- Model:
  - choose a parametric model  $P(x|y;\theta)$
  - adapt parameters  $\theta$  to data
  - How can we choose  $\theta$  to best approximate the true density  $p(x)$

# Supervised vs. Unsupervised Settings

- Task: estimate parameters
- supervised learning problems
  - given  $n$  input-output pairs  $(x_1, y_1), \dots, (x_n, y_n)$
  - $x \in \mathcal{X}$  and  $y \in \mathcal{Y}$
  - maximum likelihood (ML)
- unsupervised learning problems
  - only  $n$  observations are given:  $x_1, x_2, \dots, x_n \in \mathcal{X}$
  - (later in this lecture)

# Maximum Likelihood

- For points generated independently and identically distributed (iid) from  $p(X = x|Y = y)$ , the likelihood of the data is

$$\mathcal{L}(\theta) = \prod_{i=1}^N p(x_i|y; \theta)$$

$$\mathcal{D} = \{(x_n, y_n)\}_{n=1 \dots N}$$
$$p(\mathcal{D}|\theta)$$

- Often convenient to take logs,

$$L(\theta) = \log \mathcal{L}(\theta) = \sum_{i=1}^n \log p(x_i|y; \theta)$$

- Maximum likelihood chooses  $\theta$  to maximize  $\mathcal{L}$  (and thus  $L$ )

# Example: multinomial distribution

- Consider an experiment with  $n$  independent trials
- Each trial can result in any of  $r$  possible outcomes (e.g., a die)
- $p_i$  denotes the probability of outcome  $i$ ,  $\sum_{i=1}^r p_i = 1$
- $n_i$  denotes the number of trials resulting in outcome  $i$ ,  $\sum_{i=1}^r n_i = n$
- The likelihood is given by

$$\mathcal{L}(p_1, \dots, p_r) = \prod_{i=1}^r p_i^{n_i}$$

→ Max likelihood

- Show that the maximum likelihood estimate for  $p_i$  is  $\hat{p}_i = \frac{n_i}{n}$   
– proof in Davis & Jones, ML Estimation for the Multinomial Distribution, Teaching Statistics 14(3), 1992

# Applications

- part-of-speech tagging
  - input: sentence (=observation)
  - output: sequence of part-of-speech tags (= latent variables)
- named entity recognition (NER)
  - input: sentence (=observation)
  - output: sequence of named entites (time, person, location, organization, ...)
- protein secondary structure prediction
  - input: primary structure
  - output: secondary structure

# Example: Natural Language Processing

- Part-of-speech tagging:

– input: *Curiosity kills the cat.*

– output: <noun, verb, determiner, noun>

1 2 3 4 ↙  
| | | |

N1=20  
4<sup>20</sup>

- named entity recognition (NER)

– input: *Robert Enke was born in August 1977 in Jena.*

– output: < person, person, o, o, o, date, date, o, location>

B-P I-P

B-D, I-D B-L

✓

- NER also relevant in biomedical applications: gene/protein detection

FA

# Protein Secondary Structure Prediction

- example:

$$f(x) = \omega^T \phi(x) + b$$

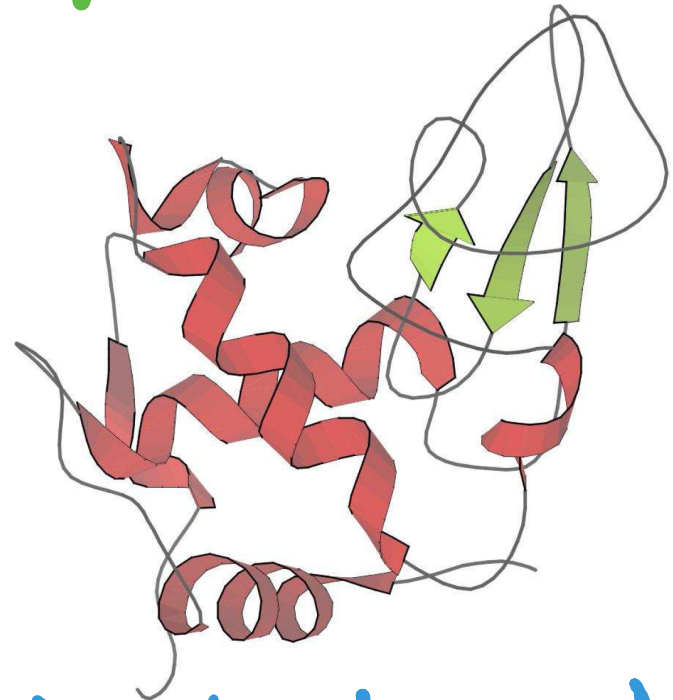

 KVFGRCELAA AMKRHGLDNY RGYSLGNWVC  
 B HHHHHH HHHHTT TTB TTB HHHHHH

AAKFESNFNT QATNRNTDGS TDYGILQINS  
 HHHHHHTTBT T EEE TTS EEETTTTEET

RWWCNDGRTP GSRNLCNIPC SALLSSDITA  
 TTB B S T T BTT SBG GGGGSSS HH

SVNCAKKIVS DGNGMNAWVA WRNRCKGTDV  
 HHHHHHHHHT SSSGGGGS HH HHHHTTTS G

QAWIRGCRL  
 GGGTTT



NLP:  $P(y_1 | k) P(y_2 | v_1, k) P(y_3 | \#_1, v_1, k) \dots$

# Naive Bayes

$$P(y|x) = \frac{P(x|y)P(y)}{P(x)}$$

$$\Rightarrow P(x_1|Y) P(x_2|Y) \dots \prod_{d=1}^D P(x_d|Y)$$



# Label Sequence Learning

- formalization:
  - input: sequence  $\xi = x_1, x_2, \dots, x_T$
  - output: sequence  $y_1, y_2, \dots, y_T$
  - elements in  $\xi$  and  $y$  are not iid!
- Structure is determined by length of input sequence

- goal:
  - prediction model:  ~~$P(\cdot|\xi)$~~
  - given a new sentence  $\xi'$ , compute prediction  $\hat{y}$ :

arg max  $P(y_1 \dots y_T | x_1 \dots x_T)$

$\gamma_1 \dots \gamma_T$

~~$P(\cdot|\xi)$~~

– capture dependencies between neighboring words

Standard

# Approaches

- ~~flat~~ approaches (naive Bayes, SVM, ...)
  - independence assumption on words of a sentence
  - cannot exploit dependencies

# Approaches

- flat approaches (naive Bayes, SVM, ...)
  - indendence assumption on words of a sentence
  - cannot exploit dependencies
- flat appraoches w/ sliding windows
  - capture dependencies within window
  - long-range dependencies are not detected

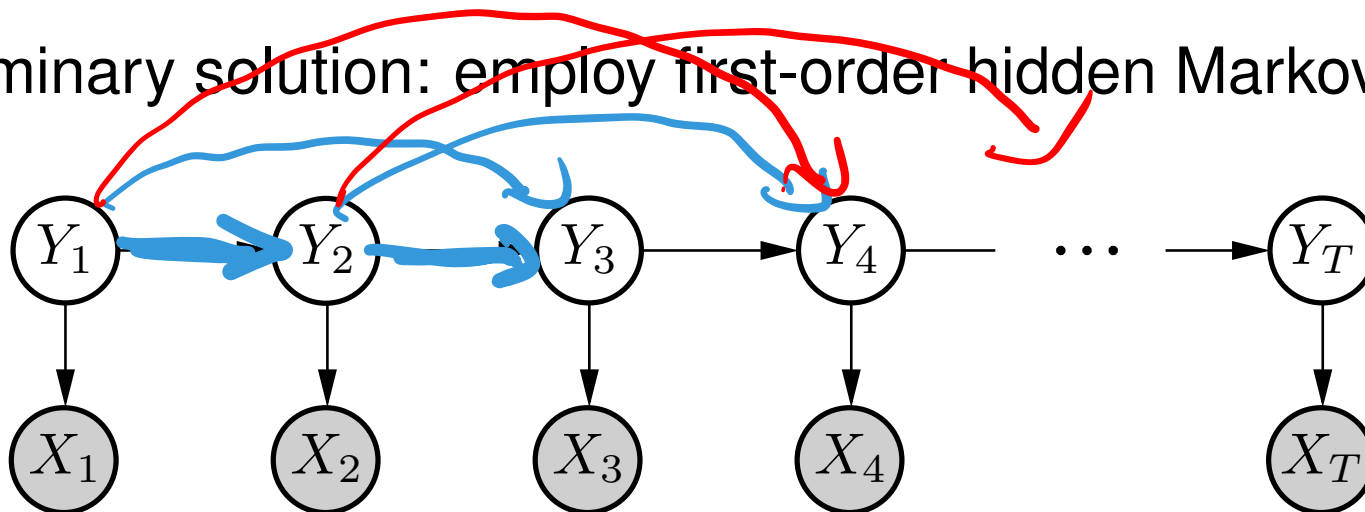
# Approaches

- flat approaches (naive Bayes, SVM, ...)
  - indendence assumption on words of a sentence
  - cannot exploit dependencies

- flat appraoches w/ sliding windows
  - capture dependencies within window
  - long-range dependencies are not detected

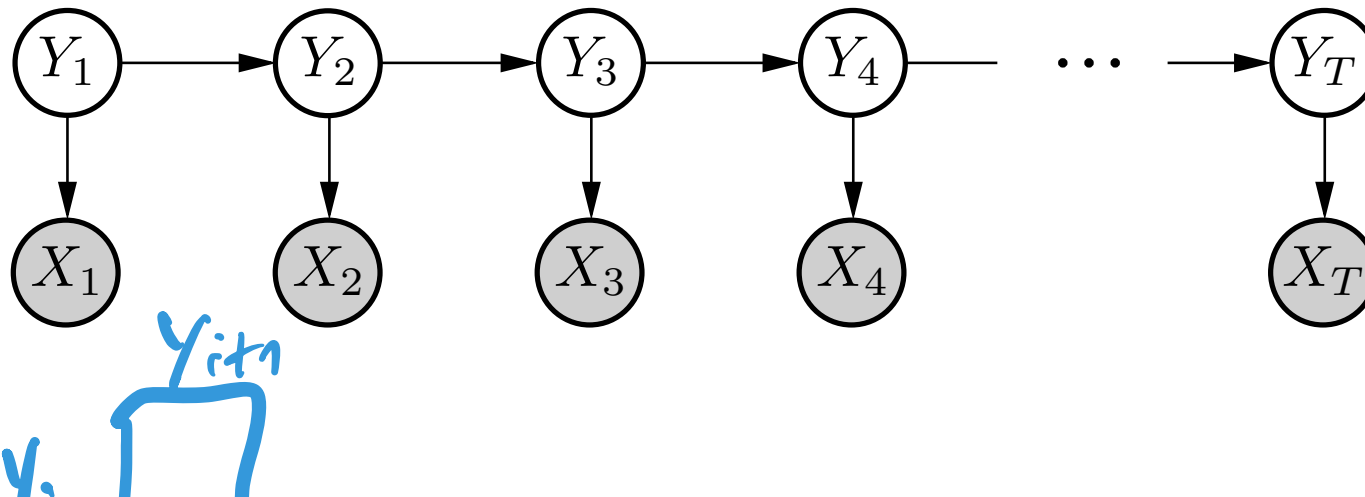
$$\begin{aligned} &P(Y_1)P(Y_2|Y_1)P(Y_3|Y_2) \\ &\dots P(X_1|Y_1) \\ &P(X_2|Y_2) \\ &\vdots \end{aligned}$$

- Preliminary solution: employ first-order hidden Markov model:

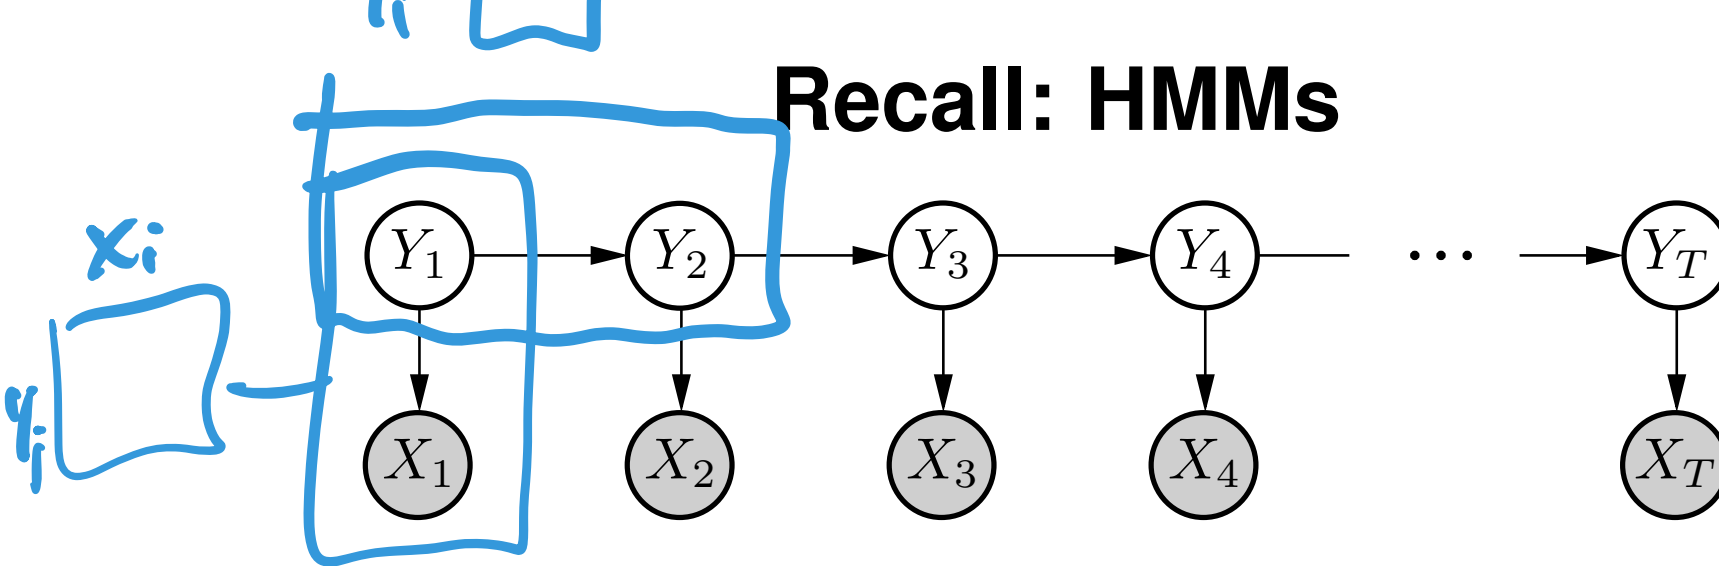


# Part-of-Speech Tagging

- Given:
  - given  $n$  pairs  $(\xi_{1,1}), \dots, (\xi_{n,n})$
  - $\xi_i = x_{i1}, \dots, x_{iT_i}$  is the  $i$ -th input sequence
  - $i = y_{i1}, \dots, y_{iT_i}$  is the  $i$ -th annotation
  - $\text{dom}(x_{ij}) = \{\text{Aachen, Aar}, \dots, \text{ZZ-top}\}$
  - $\text{dom}(y_{ij}) = \{\text{noun, verb, determiner}, \dots\}$
- Graphical model:



## Recall: HMMs



$$P(Y_1, \dots, Y_T, X_1, \dots, X_T) = P(Y_1) \left[ \prod_{t=1}^T P(Y_t | Y_{t-1}) \right] \left[ \prod_{t=1}^T P(X_t | Y_t) \right]$$

- multinomial distributions:
  - priors:  $P(Y_1)$
  - emissions:  $P(X_t | Y_t)$
  - transitions:  $P(Y_t | Y_{t-1})$

# Parameter Estimation

- Maximum likelihood says:

- Priors:  $\pi_i = P(y_1 = \sigma_i) = \frac{1}{n} \sum_{k=1}^n [[y_{k1} == \sigma_i]]$

- emissions:

$$P(x_t = w | y_t = \sigma_i) = \frac{\sum_{k=1}^n \sum_{p=1}^{T_k} [[y_{kp} == \sigma_i \wedge x_{kp} == w]]}{\sum_{k=1}^n \sum_{p=1}^{T_k} [[y_k == \sigma_i]]}$$

- transitions:

$$P(y_{t+1} = \sigma_j | y_t = \sigma_i) = \frac{\sum_{k=1}^n \sum_{p=1}^{T_k} [[y_{kp} == \sigma_i \wedge y_{k,p+1} == \sigma_j]]}{\sum_{k=1}^n \sum_{p=1}^{T_k} [[y_k == \sigma_i]]}$$

# Applying the trained HMM

- HMM can be adapted to data with maximum likelihood
- Once the probabilities are estimated, the HMM can be used for prediction
- 2 possibilities:
  - use sum-product algorithm to optimize  $P(y_t|x_1, \dots, x_T)$
  - use max-product algorithm to optimize  $P(y_1, \dots, y_T|x_1, \dots, x_T)$
  - max-product for first-order hidden Markov models is called Viterbi algorithm



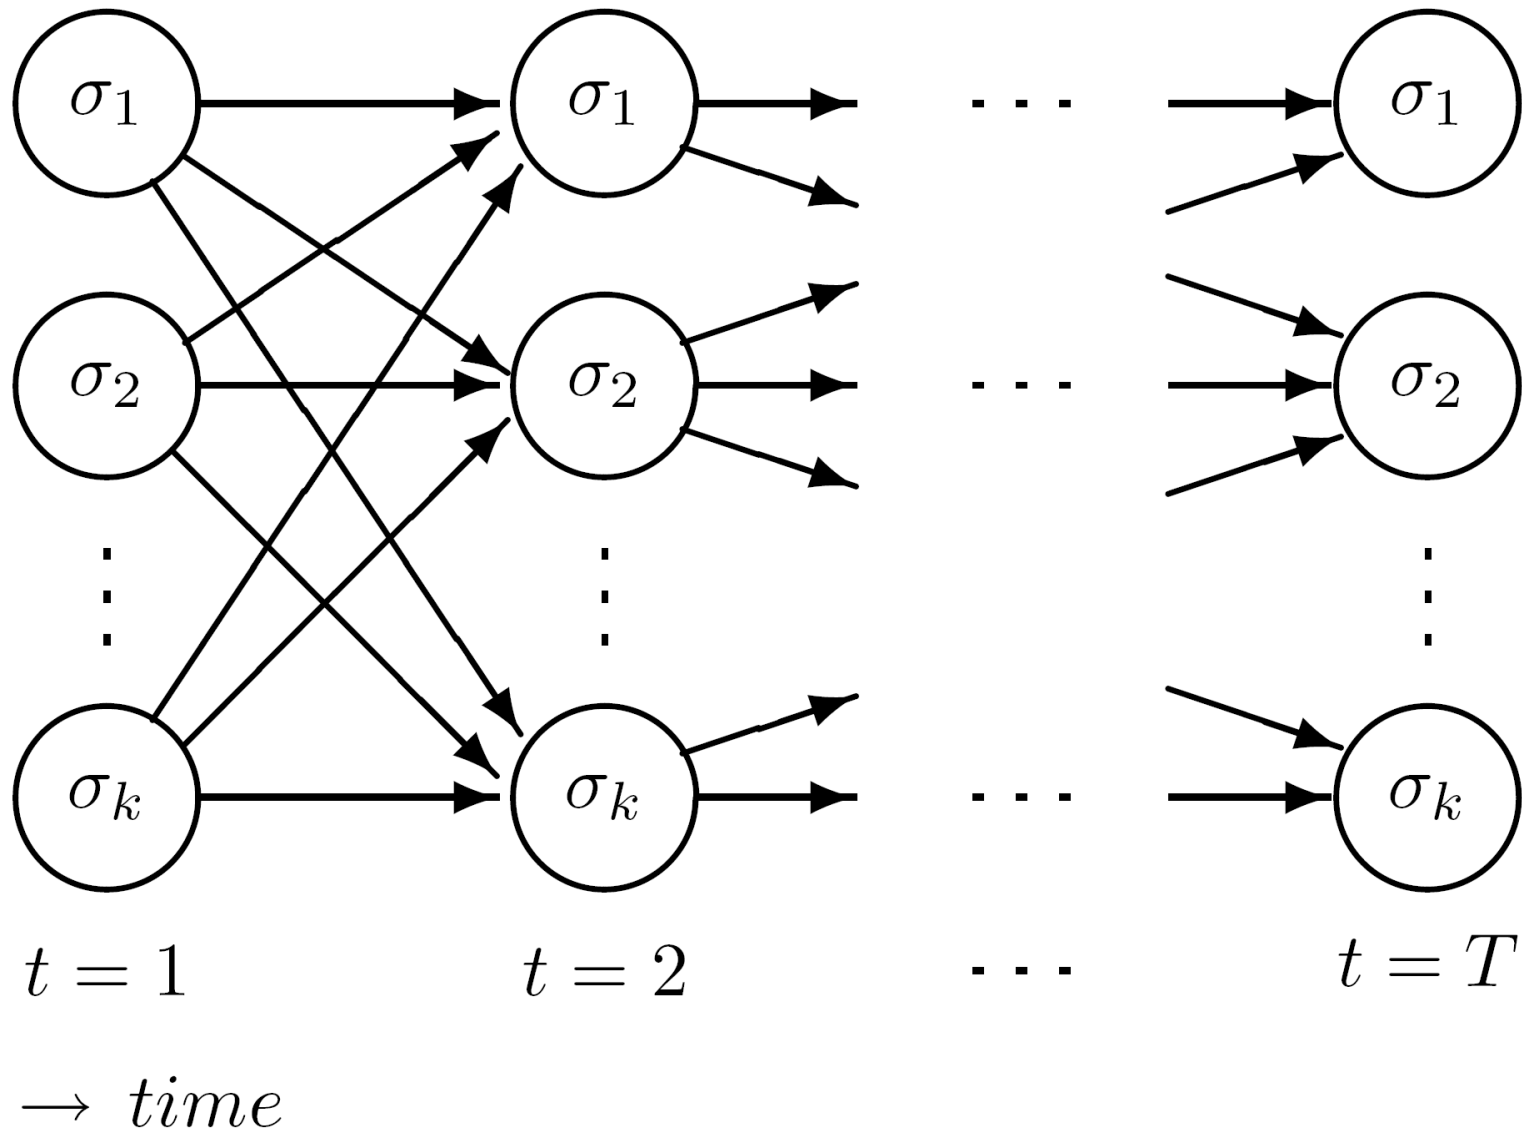
# Viterbi Algorithm

- Compute:  $_{y_1, \dots, y_T} P(y_1, \dots, y_T | x_1, \dots, x_T)$
- Define  $\delta_{t+1}(\sigma_i) = \max_{y_1, \dots, y_t} P(y_1, \dots, y_{t+1} = \sigma_i, x_1, \dots, x_{t+1})$ 
  - $\delta_{t+1}(\sigma_i)$  is the best score along a single path up to time  $t + 1$  which account for the first  $t + 1$  observations and ends in state  $\sigma_i$  at time  $t + 1$
  - apply  $\delta_{t+1}(\sigma_i)$  recursively, similar to forward-backward algorithm (except that a max than sum operation is used)
  - see also: Rabiner, Proc. IEEE 77(2), 1989 pp. 257-285

# Viterbi Algorithm

- initialize  $\delta_1(\sigma_i) = P(y_1 = \sigma_i)P(x_1|y_1 = \sigma_i)$
- initialize  $\psi_1(\sigma_i) = 0$
- loop  $j = 1, \dots, |\Sigma|$  and  $t = 1, \dots, T - 1$ :
  - $\delta_{t+1}(\sigma_j) = \left[ \max_i \delta_t(i) P(y_{t+1} = \sigma_j | y_t = \sigma_i) \right] P(x_{t+1} | y_{t+1} = \sigma_j)$
  - $\psi_{t+1}(\sigma_j) = \left[ \arg\max_i \delta_t(i) P(y_{t+1} = \sigma_j | y_t = \sigma_i) \right] P(x_{t+1} | y_{t+1} = \sigma_j)$
- termination:  $y_T^* = \arg\max_i \delta_T(\sigma_i)$
- loop  $t = T - 1, \dots, 1$ 
  - $y_t^* = \psi_{t+1}(y_{t+1}^*)$

# Trellis



# Limitations of HMMs

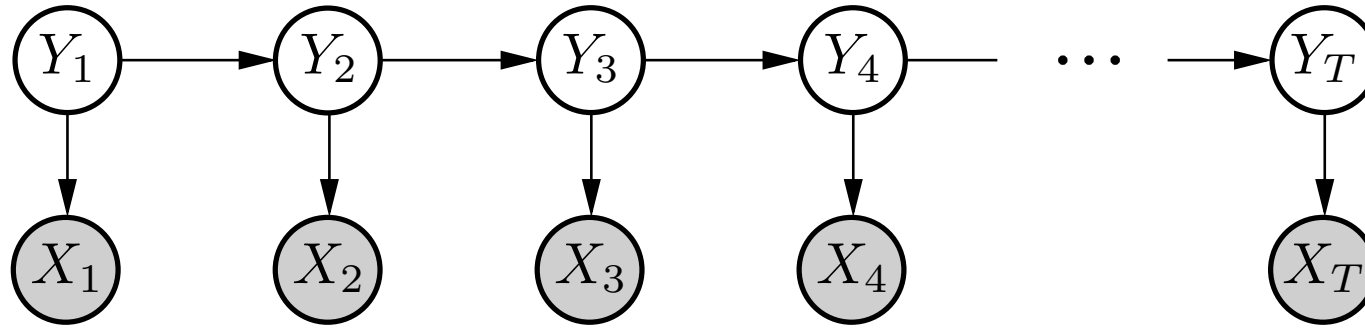
- Long-range dependencies are not captured
  - a remedy might be higher-order HMMs
  - computationally demanding
- probabilities need to be smoothed
  - unobserved words (and sequences including them) will always have zero probability
  - a common approach that does not work very well is Laplace smoothing:

$$P(x_t = w | y_t = \sigma_i) = \frac{1 + \sum_{k=1}^n \sum_{p=1}^{T_k} [[y_{kp} == \sigma_i \wedge x_{kp} == w]]}{|dom(x_t)| + \sum_{k=1}^n \sum_{p=1}^{T_k} [[y_k == \sigma_i]]}$$

# More Severe Limitations of HMMs

- HMMs are generative models
  - HMMs address the joint probability  $P(\xi, )$
  - we are interested in discriminative models  $P(|\xi)$
  - HMMs optimize the wrong criterion!
- Next time:
  - Use Markov random field instead of Bayesian network
  - Condition joint probability on the observations
  - Conditional random fields

# Recall: HMMs



- Hidden Markov models
  - generative models for sequential data
  - parameters: prior, transition, and observation probabilities
  - joint probability:

$$P(X_1, \dots, Y_1 \dots) = P(Y_1) \prod_{i=1}^T P(X_i | Y_i) \prod_{i=2}^T P(Y_i | Y_{i-1})$$

# Learning HMMs

- given:  $n$  labeled sequences  $(\xi_{1,1}), \dots, (\xi_{n,n})$
- maximum Likelihood (ML)
  - adapt parameters of HMM to data
  - HMM: ML reduces to counting
  - efficient (one pass over data suffices)
  - easy to implement
  - exact inference (Viterbi algorithm)
- drawbacks
  - $P(\text{unobserved token} | Y_i) = 0$  (remedy: smoothing techniques)
  - generative models optimize the wrong criterion

# Today: From HMMs to CRFs

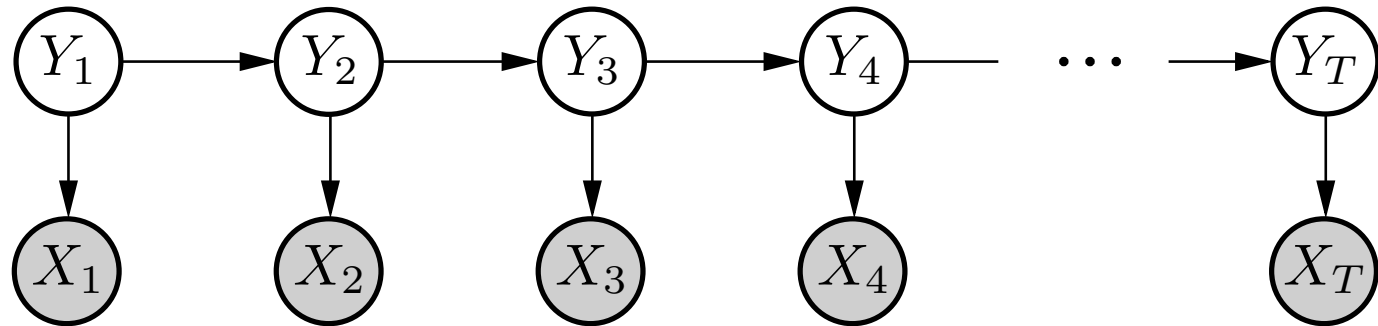
- Use undirected graphical model
  - no assumption on directions of dependencies (i.e., WWW, NLP, images, ...)
  - sequences: factor graph does not change
  - Markov random fields
- Condition joint probability of MRF on observations
  - criterion: prediction model
  - now: conditional (=discriminative) model



# Conditional Random Fields

# Markov Random Fields

HMM:



MRF:

- every BN can be translated into equivalent MRF (moralization) <sub>44/??</sub>

# MRF: Joint Probability Distribution

- joint probability factorizes across cliques
  - cliques between transitions and label-observation pairs

$$P(X_1, \dots, Y_1, \dots) = \frac{1}{Z} \prod_{i=1}^T \psi^{obs}(X_i, Y_i) \prod_{i=2}^T \psi^{trans}(Y_{i-1}, Y_i)$$

- potential functions  $\psi^{trans}(Y_i, Y_{i-1})$ ,  $\psi^{obs}(X_i, Y_i)$
- $Z$  normalization term (partition function)

# Partition Function

$$P(X_1, \dots, Y_1, \dots) = \frac{1}{Z} \prod_{i=1}^T \psi^{obs}(X_i, Y_i) \prod_{i=2}^T \psi^{trans}(Y_{i-1}, Y_i)$$

- the partition function needs to sum over all possible assignments of input and output sequences
  - we have:

$$Z = \sum_{x_1, \dots, x_T} \sum_{y_1, \dots, y_T} \prod_{i=1}^T \psi^{obs}(X_i, Y_i) \prod_{i=2}^T \psi^{trans}(Y_{i-1}, Y_i)$$

- important for  $P(X_1, \dots, Y_1, \dots)$  being a probability

# Potential Functions

$$P(X_1, \dots, Y_1, \dots) = \frac{1}{Z} \prod_{i=1}^T \psi^{obs}(X_i, Y_i) \prod_{i=2}^T \psi^{trans}(Y_{i-1}, Y_i)$$

- potential functions  $\psi^{trans}$  (transitions),  $\psi^{obs}$  (label-observ.)
  - arbitrary, non-negative, positive functions
  - capture relevant dependencies
  - defined across cliques
- problem:
  - size of largest clique depends on input (i.e., WWW)
  - remedy: represent only cliques of size 2 (=Markov network)

# Representation

$$P(X_1, \dots, Y_1, \dots) = \frac{1}{Z} \prod_{i=1}^T \psi^{obs}(X_i, Y_i) \prod_{i=2}^T \psi^{trans}(Y_{i-1}, Y_i)$$

- sequences
  - all cliques are of size 2
  - only their number varies with  $T$
- how to choose  $\psi^{trans}$ ,  $\psi^{obs}$ ?
  - (remember they have to capture relevant dependencies)
- common assumption (Hammersley & Clifford theorem):
  - $\psi$  is log-linear combination of basis functions  $\phi_j$

# Members in the Exponential Family

- basis functions:

$$\psi^{trans}(Y_i, Y_{i-1}) = \exp\left\{\sum_{j=1}^{d_{trans}} w_j^{trans} \phi_j^{trans}(Y_{i-1}, Y_i)\right\}$$

$$\psi^{obs}(X_i, Y_i) = \exp\left\{\sum_{j=1}^{d_{obs}} w_j^{obs} \phi_j^{obs}(X_i, Y_i)\right\}$$

– math turns out to be nice!

– write:

$$P(X_1, \dots, Y_1, \dots) = \frac{1}{Z} \prod_{i=1}^T \exp\left\{\sum_{j=1}^{d_{obs}} w_j \phi_j^{obs}(X_i, Y_i)\right\} \prod_{i=2}^T \exp\left\{\sum_{j=1}^{d_{trans}} w_j \phi_j^{trans}(Y_{i-1}, Y_i)\right\}$$

# Basis Functions: label-label

$$\psi^{trans}(Y_i, Y_{i-1}) = \exp \left\{ \sum_{j=1}^{d_{trans}} w_j^{trans} \phi_j^{trans}(Y_{i-1}, Y_i) \right\}$$

- simple case: indicator functions

$$\phi_1^{trans}(Y_{i-1}, Y_i) = [[Y_{i-1} = \text{noun} \wedge Y_i = \text{noun}]]$$

$$\phi_2^{trans}(Y_{i-1}, Y_i) = [[Y_{i-1} = \text{noun} \wedge Y_i = \text{verb}]]$$

$$\vdots$$
$$\vdots$$

$$\phi_{d_{trans}}^{trans}(Y_{i-1}, Y_i) = [[Y_{i-1} = \text{adverb} \wedge Y_i = \text{adverb}]]$$

- similar to HMM
- later more...



# Basis Functions: label-observation

$$\psi^{obs}(X_i, Y_i) = \exp \left\{ \sum_{j=1}^{d_{obs}} w_j^{obs} \phi_j^{obs}(X_i, Y_i) \right\}$$

- simple case: indicator functions

$$\phi_1^{obs}(X_i, Y_i) = [[X_i = \text{Aachen} \wedge Y_i = \text{noun}]]$$

$$\phi_2^{obs}(X_i, Y_i) = [[X_i = \text{Aar} \wedge Y_i = \text{noun}]]$$

$$\vdots \qquad \qquad \qquad \vdots$$

$$\phi_{d_{obs}}^{obs}(X_i, Y_i) = [[X_i = \text{ZZ-top} \wedge Y_i = \text{adverb}]]$$

- similar to HMM
- later more...

# Putting Everything Together...

$$\begin{aligned}
 P(\xi, ) &= \frac{1}{Z} \prod_{i=1}^T \exp\left\{\sum_{j=1}^{d_o} w_j^o \phi_j^o(x_i, y_i)\right\} \prod_{i=2}^T \exp\left\{\sum_{j=1}^{d_t} w_j^t \phi_j^t(y_{i-1}, y_i)\right\} \\
 &= \frac{1}{Z} \prod_{i=1}^T \exp\{\langle \wedge^o, \phi^o(x_i, y_i) \rangle\} \prod_{i=2}^T \exp\{\langle \wedge^t, \phi^t(y_{i-1}, y_i) \rangle\} \\
 &= \frac{1}{Z} \exp\left\{\sum_{i=1}^T \langle \wedge^o, \phi^o(x_i, y_i) \rangle\right\} \exp\left\{\sum_{i=2}^T \langle \wedge^t, \phi^t(y_{i-1}, y_i) \rangle\right\} \\
 &= \frac{1}{Z} \exp\left\{\langle \wedge^o, \sum_{i=1}^T \phi^o(x_i, y_i) \rangle\right\} \exp\left\{\langle \wedge^t, \sum_{i=2}^T \phi^t(y_{i-1}, y_i) \rangle\right\} \\
 &= \frac{1}{Z} \exp\left\{\left\langle \underbrace{\begin{pmatrix} \wedge^o \\ \wedge^t \end{pmatrix}}_{=:\wedge}, \underbrace{\begin{pmatrix} \sum_{i=1}^T \phi^o(x_i, y_i) \\ \sum_{i=2}^T \phi^t(y_{i-1}, y_i) \end{pmatrix}}_{=:\Phi(\xi, )} \right\rangle\right\} \\
 &= \frac{1}{Z} \exp\{\langle \wedge, \Phi(\xi, ) \rangle\}
 \end{aligned}$$

# Joint Feature Representation

- joint representation of input and output variables:

$$\Phi(\xi, ) = (\sum_{i=1}^T \phi^o(x_i, y_i), \sum_{i=2}^T \phi^t(y_{i-1}, y_i))'$$

- Example for HMM-alike basis functions:

- $\Phi(\xi, )$  counts how many times ...
- ... a *noun* is followed by *verb* (summing over transitions)
- ... the token *Aachen* is observed as a noun (sum over obs-label)
- dimensionality of  $\Phi$  is  $\text{dom}(x_i) \times \text{dom}(y_i) + \text{dom}(y_i)^2$

- POS-tagging:

- dictionary size 20,000 tokens, 36 POS-tags,  $\text{dim}(\Phi) = 721296$

# Example

# Features

- Features are engineered to capture important relations/dependencies
- all time favorites for natural language text:
  - n-grams (English: *-ing*, German: *-ung*, *-heit*, *-keit*)
  - surface clues (capitalization, all-caps, ...)
  - foreign symbols ( $\alpha$ ,  $\omega$ , ...)
  - numbers (42, 1984, ...)
- CRFs allow for rich feature spaces
  - CRFs may contain any number of basis functions
  - basis functions can be defined on the entire input sequence
  - basis functions do need not have a probabilistic interpretation.

# More Features / Relation to HMM

- observation-label/transitions can depend on input
  - $\phi^{trans}(y_{t-1}, y_t) \rightarrow \phi^{trans}(y_{t-1}, y_t; x_t)$
  - or even:  $\phi^{trans}(y_{t-1}, y_t) \rightarrow \phi^{trans}(y_{t-1}, y_t; \xi)$
  - similarly:  $\phi^{obs}(x_t, y_t) \rightarrow \phi^{obs}(\xi, y_t)$
  - (alternative graph structure)
- Implications for HMMs
  - Multi-bernoulli/nomial distribution
  - Generally infeasible

# The Exponential Family

$$P(\xi, ) = \frac{1}{Z} \exp\{\langle \wedge, \Phi(\xi, ) \rangle\}$$

- $P(\xi, )$  is a member in the exp. family, rewrite in canonical form

$$P(\xi, ) = \exp\{\langle \wedge, \Phi(\xi, ) \rangle - \log Z\}$$

- Identify the terms:
  - $\Phi(\xi, )$  is the sufficient statistics
  - $\wedge$  is the natural parameter
  - $\log Z < \infty$  is the moment generating function

# Conditional Markov Random Fields

- joint probability

$$P(\xi, ) = \frac{1}{Z} \exp\{\langle \wedge, \Phi(\xi, ) \rangle\}$$

– partition function:  $Z = \sum_{\xi} \sum_{\exp\{\langle \wedge, \Phi(\xi, ) \rangle\}}$

- condition on the observation

– apply the rule:  $P(|\xi) = P(\xi, ) / P(\xi)$

– obtain new partition function:

$$Z(\xi) = \sum_{\exp\{\langle \wedge, \Phi(\xi, ) \rangle\}}$$

- obtain a so-called conditional random field (CRF)

$$P(|\xi) = \frac{1}{Z(\xi)} \exp\{\langle \wedge, \Phi(\xi, ) \rangle\}$$



# Training CRFs with Maximum Likelihood

- given  $n$  input output examples  $(\xi_{1,1}), \dots, (\xi_{n,n})$
- the log-likelihood is given

$$\log \mathcal{L} = \sum_{i=1}^n \langle \wedge, \Phi(\xi_{i,i}) \rangle - \log Z(\wedge | \xi_i)$$

- differentiating wrt  $\wedge$  gives

$$\frac{\partial}{\partial \wedge} \log \mathcal{L} = \mathbf{E}_{\hat{p}(X,Y)}[\Phi(X,Y)] - \sum_{i=1}^n \mathbf{E}_{p(Y|\xi_i;\wedge)}[\Phi(Y,\xi_i)]$$

- empirical distribution of data  $\hat{p}$
- model distribution  $p$

# Optimization

- direct optimization is expensive and often infeasible
  - E.g., calculating the partition function is time consuming if at all possible
- Many different optimization strategies have been proposed:
  - linear programming (Roth & Li, 2005)
  - iterative scaling (Lafferty et al., 2001)
  - conjugate gradients (Sha & Pereira, 2003)
  - Gauss-Newton subspace optimization (Altun et al., 2004)
  - gradient tree boosting (Dietterich et al., 2004)
  - stochastic meta descent (Vishwanathan et al., 2006)
  - perceptron algorithm (Altun et al., 2003)
  - ...

# The Perceptron Algorithm for CRFs

# CRF vs. HMM

- characteristics:
  - CRF: undirected graph, conditional models
  - HMM: directed BN, generative model
- CRFs generalize HMMs
  - CRFs allow for rich feature spaces
  - HMMs restricted to implicit bag-of-words representation
- Optimization
  - CRF: difficult, complex optimization problem
  - HMM: simple, easy to implement
- Similarities:
  - inference algorithms (Viterbi, sum-product)

# Posterior vs. MAP

- Once optimal parameters  $\hat{\Lambda}^*$  are found these are used as plug-in estimates  $P(|\xi; \hat{\Lambda}^*)$ 
  - posterior distribution allows for computing confidence intervals
- However, the full posterior is not always needed
  - often, the maximum a posteriori (MAP) estimate suffices
  - e.g., prediction model  $\hat{P}(|\xi)$
  - computing MAP estimates is much cheaper than full posterior!

# Computing MAP Estimates

- For MAP estimates compute

$$\begin{aligned}\hat{\Lambda} &= P(\Lambda | \xi) \\ &= \frac{1}{Z(\xi)} \exp\{\langle \Lambda, \Phi(\xi, \cdot) \rangle\} \\ &= \langle \Lambda, \Phi(\xi, \cdot) \rangle\end{aligned}$$

– because  $\exp$  is a monotone function and  $\frac{1}{Z(\xi)}$  is constant

- We arrive at:

$$P(\xi, \cdot) \propto \underbrace{\langle \Lambda, \Phi(\xi, \cdot) \rangle}_{=: f(\xi, \cdot)}$$

# Outlook

- adapt  $f(\xi, ) = \langle \wedge, \Phi(\xi, ) \rangle$  to data
- perceptron algorithm
  - primal: efficient, nof parameters =  $\dim(\Phi)$
  - dual: nof parameters = nof possible output sequences
- dual perceptron
  - explicit representation is infeasible
  - solve implicitly by column generation
- examples

# Recall: Sequential CRFs

$$P(y|\xi) = \frac{1}{Z(\xi)} \prod_{i=1}^T \psi^{obs}(X_i, Y_i) \prod_{i=2}^T \psi^{trans}(Y_{i-1}, Y_i)$$

- potential functions:

$$\psi^{trans}(Y_i, Y_{i-1}) = \exp \left\{ \sum_{j=1}^{d_{trans}} w_j^{trans} \phi_j^{trans}(Y_{i-1}, Y_i) \right\}$$

$$\psi^{obs}(X_i, Y_i) = \exp \left\{ \sum_{j=1}^{d_{obs}} w_j^{obs} \phi_j^{obs}(X_i, Y_i) \right\}$$

- $Z(\xi) = \sum_{y_1, \dots, y_T} \prod_{i=1}^T \psi^{obs}(X_i, Y_i) \prod_{i=2}^T \psi^{trans}(Y_{i-1}, Y_i)$



# Exemplary Basis Functions

- label-label indicator functions:

$$\phi_1^{trans}(Y_{i-1}, Y_i) = [[Y_{i-1} = \text{noun} \wedge Y_i = \text{noun}]]$$

$$\phi_2^{trans}(Y_{i-1}, Y_i) = [[Y_{i-1} = \text{noun} \wedge Y_i = \text{verb}]]$$

$$\vdots$$
$$\vdots$$

$$\phi_{d_{trans}}^{trans}(Y_{i-1}, Y_i) = [[Y_{i-1} = \text{adverb} \wedge Y_i = \text{adverb}]]$$

- label-observation indicators:

$$\phi_1^{obs}(X_i, Y_i) = [[X_i = \text{Aachen} \wedge Y_i = \text{noun}]]$$

$$\phi_2^{obs}(X_i, Y_i) = [[X_i = \text{Aar} \wedge Y_i = \text{noun}]]$$

$$\vdots$$
$$\vdots$$

$$\phi_{d_{obs}}^{obs}(X_i, Y_i) = [[X_i = \text{ZZ-top} \wedge Y_i = \text{adverb}]]$$

# Joint Feature Representation

- Joint representation of input and output variables:

$$\Phi(\xi, ) = \left( \sum_{i=1}^T \phi^o(x_i, y_i)', \sum_{i=2}^T \phi^t(y_{i-1}, y_i)' \right)'$$

- Rewrite conditional probability:

$$P(|\xi) = \frac{1}{Z(\xi)} \exp \left\{ \langle \wedge, \Phi(\xi, ) \rangle \right\}$$

- Observation:

$$P(|\xi) \propto \langle \wedge, \Phi(\xi, ) \rangle$$

- MAP estimate:

$$\hat{\cdot} = P(\cdot|\xi) = \langle \wedge, \Phi(\xi, \cdot) \rangle$$

# Example

- $\xi = \text{Bob jagt den Hund}$
- We want

$$[N, V, A, N] = \langle \wedge, \Phi(\xi, \cdot) \rangle$$

# Example

- $\xi = \text{Bob jagt den Hund}$
- We want

$$[N, V, A, N] = \langle \wedge, \Phi(\xi, \vec{\cdot}) \rangle$$

- Equivalent representation:

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, A, A]) \rangle$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, A, N]) \rangle$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, N, A]) \rangle$$

$$\vdots > \vdots$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [V, V, V, V]) \rangle$$

# Example Contd.

- Another equivalent representation:

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle - \langle \wedge, \Phi(\xi, [A, A, A, A]) \rangle > 0$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle - \langle \wedge, \Phi(\xi, [A, A, A, N]) \rangle > 0$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle - \langle \wedge, \Phi(\xi, [A, A, N, A]) \rangle > 0$$

$$\vdots \qquad \qquad \vdots$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle - \langle \wedge, \Phi(\xi, [V, V, V, V]) \rangle > 0$$

- The other way round:
  - Update weight vector  $\wedge$  in case of an error:

$$\langle \wedge, \Phi(\xi_{i,i}) \rangle - \max \langle \wedge, \Phi(\xi_i, \vec{\cdot}) \rangle < 0$$

# Primal Perceptron

- Simplify things:
  - Error:  $i \neq \hat{y} = \langle \wedge, \Phi(\xi_i, \cdot) \rangle$

- Recall gradient of CRF:

$$\frac{\partial \log \mathcal{L}}{\partial \wedge} = \underbrace{\mathbf{E}_{\hat{p}(X,Y)}[\Phi(X,Y)]}_{\text{truth/emp. distr.}} - \underbrace{\sum_{i=1}^n \mathbf{E}_{p(Y|\xi_i;\wedge)}[\Phi(Y,\xi_i)]}_{\text{prediction of model/model distr.}}$$

- Perceptron: perform gradient steps if  $i$ -th example is incorrect:

$$\wedge \leftarrow \wedge + \underbrace{\Phi(\xi_{i,i})}_{\text{true pair}} - \underbrace{\Phi(\xi_i, \hat{\cdot})}_{\text{erroneous prediction}}$$

# Primal Perceptron Algorithm

```
1 loop  $r = 1, \dots, r_{max}$ 
2   loop  $i = 1, \dots, n$ 
3     Compute  $\hat{y} = \langle \wedge, \Phi(\xi_i, \cdot) \rangle$ 
4     If  $y_i \neq \hat{y}$ 
5       Update  $\wedge \leftarrow \wedge + \Phi(\xi_{i,i}) - \Phi(\xi_i, \cdot)$ 
6     End (if)
7   End loop ( $i$ )
8 End loop ( $r$ )
```

# Convergence

## Theorem (Extension of Novikoff)

Given  $n$  labeled examples  $(\xi_{1,1}), \dots, (\xi_{n,n})$ , with  $i \in \mathcal{Y}(\xi_i)$ . Let  $r$  be the radius of the smallest hypersphere enclosing all difference vectors  $\Phi(\xi_{i,i}) - \Phi(\xi_i, \neg)$ , for all  $i$  and  $\neg \neq i$ ,

$$r = \max_{1 \leq i \leq n} \max_{\substack{\neg \in \mathcal{Y}(\xi_i) \\ \neg \neq i}} |\Phi(\xi_{i,i}) - \Phi(\xi_i, \neg)|.$$

If there exists a vector  $\wedge^*$  such that

$$\forall_{i=1}^n \forall_{\neg \in \mathcal{Y}(\xi_i)} \langle \wedge^*, \Phi(\xi_{i,i}) \rangle - \langle \wedge^*, \Phi(\xi_i, \neg) \rangle \geq \bar{\gamma} \quad (2)$$

holds for some  $\bar{\gamma} > 0$  then the number of update steps of the generalized perceptron algorithm is upper bounded by

$$\left( \frac{r}{\bar{\gamma}} \right)^2 |\wedge|^2. \quad (3)$$



# Proof of Theorem

Proof. The weight vector is initialized with  $\wedge^{(0)} = \mathbf{0}$ . Let  $t > 0$  indicate the  $t$ -th error of the generalized perceptron, that is for some  $1 \leq i \leq n$

$$i \neq \hat{i} = \arg \min_{j \in \mathcal{Y}(\xi_i)} \langle \wedge^{(t-1)}, \Phi(\xi_i, j) \rangle.$$

The corresponding update step is given by

$$\wedge^{(t)} = \wedge^{(t-1)} + \Phi(\xi_{i,i}) - \Phi(\xi_i, \hat{i}) \quad (4)$$

Multiplying Equation 4 with the optimal weight vector  $\wedge^*$  yields

$$\begin{aligned} \langle \wedge^*, \wedge^{(t)} \rangle &= \langle \wedge^*, \wedge^{(t-1)} \rangle + \langle \wedge^*, \Phi(\xi_{i,i}) \rangle - \langle \wedge^*, \Phi(\xi_i, \hat{i}) \rangle \\ &\geq \langle \wedge^*, \wedge^{(t-1)} \rangle + \bar{\gamma} \end{aligned}$$

Applying the principle of induction gives us  $\langle \wedge^*, \wedge^{(t)} \rangle \geq t\bar{\gamma}$ .

# Proof of Theorem (Contd.)

Now we bound  $|\wedge^{(t)}|^2$  from above by

$$\begin{aligned} |\wedge^{(t)}|^2 &= \langle \wedge^{(t-1)} + \Phi(\xi_{i,i}) - \Phi(\xi_{i,\hat{i}}), \wedge^{(t-1)} + \Phi(\xi_{i,i}) - \Phi(\xi_{i,\hat{i}}) \rangle \\ &= |\wedge^{(t-1)}|^2 + 2\langle \wedge^{(t-1)}, \Phi(\xi_{i,i}) - \Phi(\xi_{i,\hat{i}}) \rangle + |\Phi(\xi_{i,i}) - \Phi(\xi_{i,\hat{i}})|^2 \\ &\leq |\wedge^{(t-1)}|^2 + |\Phi(\xi_{i,i}) - \Phi(\xi_{i,\hat{i}})|^2 \\ &\leq |\wedge^{(t-1)}|^2 + r^2. \end{aligned}$$

Thus, by induction we have  $|\wedge^{(t)}|^2 \leq tr^2$ . Putting everything together gives us

$$\begin{aligned} t\bar{\gamma} &\leq \langle \wedge^*, \wedge^{(t)} \rangle \\ &\leq |\wedge^*| |\wedge^{(t)}| \\ &\leq |\wedge^*| \sqrt{tr}. \end{aligned}$$

Solving for  $t$  implies the upper bound

$$t \leq \left( \frac{r}{\bar{\gamma}} \right)^2 |\wedge^{(*)}|^2.$$

# Towards Dual Perceptrons

- Observation:  $\wedge^{(0)} \leftarrow \mathbf{0}$
- $i$ -th example violates constraint:
  - Update:  $\wedge^{(t+1)} = \wedge^{(i)} + \langle \wedge, \Phi(\xi_i, i) \rangle - \max \langle \wedge, \Phi(\xi_i, \cdot) \rangle$
- Idea: remember how many times the pair  $(\xi_i, \cdot)$  is used for an update!
  - Variable  $\alpha_i(\cdot)$  acts as a counter
  - Initialize:  $\alpha_i(\cdot) \leftarrow 0$
  - Update:  $\alpha_i(\cdot) \leftarrow \alpha_i(\cdot) + 1$
- The  $\alpha$  are bound to violated constraints!

# Dual Representation

- Dual parameters
  - $\alpha_i(\vec{\gamma})$  is proportional to the importance of  $\Phi(\xi_{i,i}) - \Phi(\xi_i, \vec{\gamma})$
- Recall:  $\Lambda$  counted the number of updates for  $\Lambda$ 
  - we can thus write:

$$\Lambda = \sum_{i=1}^n \sum_{\vec{\gamma} \neq \vec{\gamma}_i} \alpha_i(\vec{\gamma}) (\Phi(\xi_{i,i}) - \Phi(\xi_i, \vec{\gamma}))$$

- Sparse representation
  - Generally, there are exponentially many  $\vec{\gamma} \neq \vec{\gamma}_i$
  - However, only a few of them will have an  $\alpha_i(\vec{\gamma}) > 0$
  - Feature vector can efficiently be encoded and stored (compare dimensionality of primal and dual!)

# Dual Decision Function

$$\Lambda = \sum_{i=1}^n \sum_{\neg \neq_i} \alpha_i(\neg) (\Phi(\xi_{i,i}) - \Phi(\xi_i, \neg))$$

- Plug dual representation of  $\Lambda$  into decision function:

$$\begin{aligned} f(\xi',') &= \langle \Lambda, \Phi(\xi',') \rangle \\ &= \left\langle \sum_{i=1}^n \sum_{\neg \neq_i} \alpha_i(\neg) (\Phi(\xi_{i,i}) - \Phi(\xi_i, \neg)), \Phi(\xi',') \right\rangle \\ &= \sum_{i=1}^n \sum_{\neg \neq_i} \alpha_i(\neg) (\langle \Phi(\xi_{i,i}) - \Phi(\xi_i, \neg), \Phi(\xi',') \rangle) \\ &= \sum_{i=1}^n \sum_{\neg \neq_i} \alpha_i(\neg) (\langle \Phi(\xi_{i,i}), \Phi(\xi',') \rangle - \langle \Phi(\xi_i, \neg), \Phi(\xi',') \rangle) \end{aligned}$$

# Kernels and the Dual Perceptron

- Define  $K(\xi, \cdot, \xi', \cdot) = \langle \Phi(\xi, \cdot), \Phi(\xi', \cdot) \rangle$ 
  - $K$  is called kernel
  - computes inner product in space spanned by  $\Phi$
  - rewrite  $f(\xi, \cdot)$  in terms of kernel functions:

$$f_D(\xi', \cdot) = \sum_{i=1}^n \sum_{\bar{\cdot} \neq i} \alpha_i(\bar{\cdot}) (K(\xi_{i,i}, \xi', \cdot) - K(\xi_{i,\bar{\cdot}}, \xi', \cdot))$$

- Example (sequences, indicator functions)

$$\begin{aligned} K(\xi, \cdot, \bar{\xi}, \bar{\cdot}) &= \langle \Phi(\xi, \cdot), \Phi(\bar{\xi}, \bar{\cdot}) \rangle \\ &= \sum_{s,t} [[y^{s-1} = \bar{y}^{t-1} \wedge y^s = \bar{y}^t]] \\ &\quad + \sum_{s,t} [[y^s = \bar{y}^t]] K_x(x^s, \bar{x}^t) \end{aligned}$$

# Kernels on Tokens

- Kernel  $K_x$  computes similarity of two tokens
  - Simplest case:  $K_x(x, x') = [[x == x']]$
  - No generalization!
- A better choice:
  - $K_x$  computes similarity of feature vectors of observations
  - e.g.,  $n$ -grams, surface clues
  - Let  $\psi(x)$  be the feature vector of token  $x$ , then

$$K_x(x, x') = \langle \psi(x), \psi(x') \rangle$$

- $K_x$  can be precomputed for the training process

# Dual Perceptron Algorithm

```
1 loop  $r = 1, \dots, r_{max}$ 
2   loop  $i = 1, \dots, n$ 
3     Compute  $\hat{y} = f_D(\xi_i, \vec{w})$ 
4     If  $y_i \neq \hat{y}$ 
5       Increment  $\alpha_i(\hat{y}) \leftarrow \alpha_i(\hat{y}) + 1$ 
6     End (if)
7   End loop ( $i$ )
8 End loop ( $r$ )
```

- Convergence
  - see Collins (2002) and Altun et al. (2003)



# What about the Argmax?

- For dual perceptron it's easy!
- Decompose  $f(\xi, ) = f_1(\xi, ) + f_2(\xi, )$  with

$$f_1(\xi, ) = \sum_{\sigma, \tau} a(\sigma, \tau) \sum_s [[y^{s-1} = \sigma \wedge y^s = \tau]]$$

$$a(\sigma, \tau) = \sum_{i, \neq_i} \alpha_i(\cdot) \sum_t [[\bar{y}^{t-1} = \sigma \wedge \bar{y}^t = \tau]]$$

- and

$$f_2(\xi, ) = \sum_{s, \sigma} [[y^s = \sigma]] \sum_{i, t} b(i, t, \sigma) K_x(x^s, x_i^t),$$

$$b(i, t, \sigma) = \sum_{\neq_i} [[y^t = \sigma]] \alpha_i(\cdot)$$

- (homework: show that  $f = f_1 + f_2$ !)

# Correspondence to Viterbi Algorithm

- $a(\sigma, \tau)$  corresponds to transition probabilities  $P(y_t = \tau | y_{t-1} = \sigma)$
- for observation scores compute:
  - $B_i^{s\sigma} = \sum_j \sum_t b(j, t, \sigma) k(x_i^s, x_j^t)$
  - $B_i^{s\sigma}$  corresponds to  $P(x_{i,s} | y_s = \sigma)$
- Note that  $a$  and  $b$  (or  $B$ ) are scores and can be interpreted as log-probs.
- $a$  and  $B$  can be directly plugged into log-Viterbi algorithm
- Equivalence between log-Viterbi ( $\log(P(|\xi))$ ) and  $f(\xi, )$

# Named Entity Recognition

**Example:** Como (O) contrapartida (O) Deutsche (C-B) Telekom (C-I) vender (O) al (O) consorcio (O) francs (O) su (O) participacion (O) del (O) por (O) ciento (O) en (O) el (O) empresa (O) mixta (O) britnica (O) MetroHoldings (C-B).

(see Altun et al. (2003))

# Natural Language Parsing

(see Collins&Duffy, 2002)

# BioCreative

- Detection of gene and protein names in biomedical abstracts

(Brefeld et al., 2005)

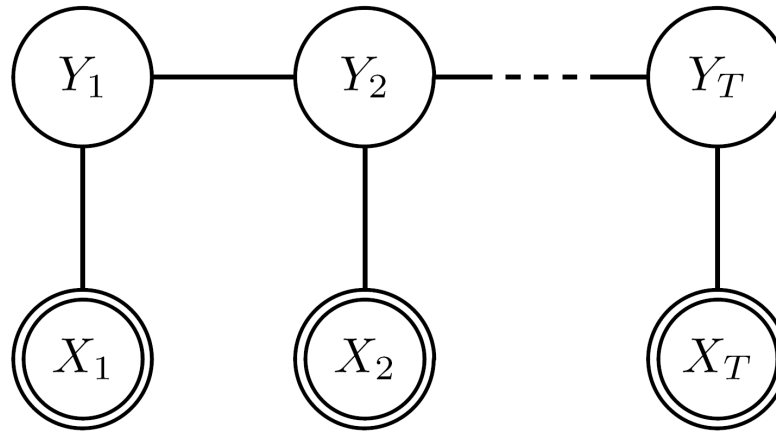
# Summary

- Perceptrons for CRFs
  - aka generalized/structured perceptron
- pos:
  - easy to implement
  - efficient training process
- neg:
  - depends on ordering
  - no confidences
  - only 0/1 loss

# Outlook

- Remedy: structural SVMs!

# Recall: Conditional MRF



$$P(\cdot|\xi) = \frac{1}{Z(\xi)} \prod_{i=1}^T \psi^{obs}(X_i, Y_i) \prod_{i=2}^T \psi^{trans}(Y_{i-1}, Y_i)$$
$$\propto \langle \wedge, \Phi(\xi, \cdot) \rangle$$



# Recall: Generalized Linear Models

- $\xi = \text{Bob jagt den Hund}$
- We want

$$[N, V, A, N] = \langle \wedge, \Phi(\xi, \cdot) \rangle$$

- Equivalent representation:

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, A, A]) \rangle$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, A, N]) \rangle$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, N, A]) \rangle$$

$$\vdots > \vdots$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [V, V, V, V]) \rangle$$

# Recall: Primal/Dual Perceptron

- Primal perceptron:

- Decision function:  $f(\xi, ) = \langle \wedge, \Phi(\xi, ) \rangle$
- Update rule:  $\wedge \leftarrow \wedge + \Phi(\xi_{i,i}) - \Phi(\xi_i, \hat{)}$

- Dual perceptron:

- Use relation:  $\wedge = \sum_{i=1}^n \sum_{\neg \neq i} \alpha_i(\neg) (\Phi(\xi_{i,i}) - \Phi(\xi_i, \neg))$
- Decision function:

$$f(\xi', ' ) = \sum_{i=1}^n \sum_{\neg \neq i} \alpha_i(\neg) (\langle \Phi(\xi_{i,i}), \Phi(\xi', ' ) \rangle - \langle \Phi(\xi_i, \neg), \Phi(\xi', ' ) \rangle)$$

- Update rule:  $\alpha_i(\neg) \leftarrow \alpha_i(\neg) + 1$

# Primal/Dual Algorithm

```
1 loop  $r = 1, \dots, r_{max}$ 
2   loop  $i = 1, \dots, n$ 
3     Compute prediction  $\hat{\cdot}$ 
4     If  $i \neq \hat{\cdot}$ 
5       Update  $\wedge$  (primal) or  $\alpha_i(\hat{\cdot})$  (dual)
6     End (if)
7   End loop ( $i$ )
8 End loop ( $r$ )
```

# How to Compute the Prediction in Step 3?

- What is the relation between...
  - Viterbi algorithm
  - max-product algorithm
  - max-sum algorithm
  - scoring function  $f(\xi, )$
  - ???
- How can we compute  $-f(\xi, \cdot)$ ?
- Answer: use log-Viterbi = max-sum algorithm!

# Recall: Homework

- Dual perceptron:
- Decompose  $f(\xi, ) = f_1(\xi, ) + f_2(\xi, )$  with

$$f_1(\xi, ) = \sum_{\sigma, \tau} a(\sigma, \tau) \sum_s [[y^{s-1} = \sigma \wedge y^s = \tau]]$$

$$a(\sigma, \tau) = \sum_{i, \neq_i} \alpha_i() \sum_t [[\bar{y}^{t-1} = \sigma \wedge \bar{y}^t = \tau]]$$

- and

$$f_2(\xi, ) = \sum_{s, \sigma} [[y^s = \sigma]] \sum_{i, t} b(i, t, \sigma) K_x(x^s, x_i^t),$$

$$b(i, t, \sigma) = \sum_{\neq_i} [[y^t = \sigma]] \alpha_i()$$

# Recall: Viterbi Algorithm

- Computes:  $_{y_1, \dots, y_T} P(y_1, \dots, y_T | x_1, \dots, x_T)$
- Viterbi = max-product algorithm
  - define  $\delta_{t+1}(\sigma) = \max_{y_1, \dots, y_t} P(y_1, \dots, y_{t+1} = \sigma, x_1, \dots, x_{t+1})$
  - best score along a single path that ends in state  $\sigma$  at time  $t + 1$
- log-Viterbi = max-sum algorithm
  - define  $\delta_{t+1}(\sigma) = \max_{y_1, \dots, y_t} \log P(y_1, \dots, y_{t+1} = \sigma, x_1, \dots, x_{t+1})$
  - apply  $\delta_{t+1}(\sigma_i)$  recursively

# Log-Viterbi Algorithm

- initialize  $\delta_1(\sigma) = \log P(y_1 = \sigma) + \log P(x_1|y_1 = \sigma)$
- initialize  $\psi_1(\sigma) = 0$
- loop  $\sigma \in \Sigma$  and  $t = 2, \dots, T$ :
  - $\delta_t(\sigma) = \left[ \max_{\tau} \delta_{t-1}(\tau) + \log P(y_t = \sigma|y_{t-1} = \tau) \right] + \log P(x_t|y_t = \sigma)$
  - $\psi_t(\sigma) = \left[ \underset{\tau}{\delta_{t-1}(\tau)} + \log P(y_t = \sigma|y_{t-1} = \tau) \right] + \log P(x_t|y_t = \sigma)$
- termination:  $y_T^* =_{\sigma} \delta_T(\sigma)$
- loop  $t = T, \dots, 2$ 
  - $y_{t-1}^* = \psi_t(y_t^*)$

# Scoring Function $f(\xi, )$

- Capture  $\log P(y_1 = \sigma)$  implicitly by adding constant label  $y_0$ .
- Observation probabilities:

$$\log P(x_t | y_t = \sigma) \propto \underbrace{\sum_j \sum_{s=1}^{T_j} \sum_{\bar{\neq}_i} [[y^t = \sigma]] \alpha_i(\bar{\cdot}) k(x_t, x_{j,s})}_{b(\sigma, x_t)}$$

- Transition probabilities:

$$\log P(y_t = \tau | y_{t-1} = \sigma) \propto \underbrace{\sum_{\bar{i}, \bar{\neq}_i} \alpha_i(\bar{\cdot}) \sum_t [[\bar{y}^{t-1} = \sigma \wedge \bar{y}^t = \tau]]}_{a(\sigma, \tau)}$$



# It holds...

## Theorem

*Given  $n$  input-output pairs of sequences of length  $T_i$  for  $1 \leq i \leq n$ , let  $\Sigma$  denote the output alphabet with  $|\Sigma| < \infty$ . Let  $f$  be defined as*

$$f(\xi, \tau) = \sum_{i=1}^n \sum_i \alpha_i(\tau) (\langle \Phi(\xi_i, i), \Phi(\xi, \tau) \rangle - \langle \Phi(\xi_i, \tau), \Phi(\xi, \tau) \rangle),$$

*where  $\Phi(\xi, \tau)$  denotes the joint feature map. Then for all  $\alpha_i(\tau) \geq 0$  and any observation sequence  $\xi$  of length  $T$ ,*

$$\hat{\tau} = \arg \max_{\tau \in \Sigma^T} f(\xi, \tau)$$

*can be computed with a Viterbi algorithm in time  $\mathcal{O}(T|\Sigma|^2)$ .*

# Proof:

- The model  $f$  has the form

$$\begin{aligned}
 f(\xi, ) &= \sum_{i=1}^n \sum_i \alpha_i(\cdot) (\langle \Phi(\xi_{i,i}), \Phi(\xi, ) \rangle - \langle \Phi(\xi_i, \cdot), \Phi(\xi, ) \rangle) \\
 &= \sum_{i=1}^n \sum_i \alpha_i(\cdot) \left( \sum_{s,t} ([y_{i,s} = y_t] - [\bar{y}_s = y_t]) k(x_{i,s}, x_t) \right. \\
 &\quad \left. + \sum_{s,t} [[y_{i,s-1} = y_{t-1} \wedge y_{i,s} = y_t] - [\bar{y}_{s-1} = y_{t-1} \wedge \bar{y}_s = y_t]] \right).
 \end{aligned}$$

- Make the dependency on labels  $\sigma, \tau \in \Sigma$  explicit by summing over all transitions and observation states

$$\begin{aligned}
 f(\xi, ) &= \sum_{\sigma, \tau \in \Sigma} \sum_{i, \bar{i} \neq i} \alpha_i(\cdot) \left( \sum_{s,t} ([y_{i,s} = \sigma] - [\bar{y}_s = \sigma]) [[y_t = \tau]] k(x_{i,s}, x_t) \right. \\
 &\quad \left. + \sum_{s,t} ([y_{i,s-1} = \sigma \wedge y_{i,s} = \tau] - [\bar{y}_{s-1} = \sigma \wedge \bar{y}_s = \tau]) \right. \\
 &\quad \left. \times [[y_{t-1} = \sigma \wedge y_t = \tau]] \right).
 \end{aligned}$$

# Proof Contd.

The transition scores from label  $\sigma$  to label  $\tau$  are now given by

$$a(\sigma, \tau) = \sum_{i=1}^n \sum_i \alpha_i(\cdot) \left( \sum_{t=1}^{T_i} [[y_{i,t-1} = \sigma \wedge y_{i,t} = \tau]] - [[\bar{y}_{t-1} = \sigma \wedge \bar{y}_t = \tau]] \right)$$

and observation scores for label  $y_s = \sigma$  and observation  $x_s$  by

$$b(\sigma, x) = \sum_{i=1}^n \sum_{t=1}^{T_i} \sum_i \alpha_i(\cdot) ([y_{i,t} = \sigma] - [\bar{y}_t = \sigma]) k(x_{i,t}, x).$$

...

The hypothesis  $f(\xi, )$  can be rewritten in terms of transition scores  $a(\sigma, \tau)$  and observation scores  $b(\sigma, x)$

$$f(\xi, ) = \underbrace{\sum_{\sigma, \tau \in \Sigma} a(\sigma, \tau) \sum_{s=1}^T [[y_{s-1} = \sigma \wedge y_s = \tau]]}_{=: f_a(\xi, )} + \underbrace{\sum_{s=1}^T \sum_{\sigma \in \Sigma} [[y_s = \sigma]] b(\sigma, x_s)}_{=: f_b(\xi, )}.$$

where  $f_a$  weights the occurrences of neighboring labels in by corresponding scores of the model and  $f_b$  determines how well observations  $x_s$  fit to their labels  $y_s$  given the model. To decode the top scoring sequence we define

$$\delta_t(\sigma) = \max_{y_1, \dots, y_{t-1}} f(\xi, y_1, \dots, y_{t-1}, y_t = \sigma), \quad (5)$$

that is,  $\delta_t(\sigma)$  denotes the top scoring partial sequence up to position  $t - 1$  where  $y_t = \sigma$ .

# Mathematical Induction: The Base Case

We first show by induction that

$$\delta_{t+1}(\sigma) = \max_{\tau \in \Sigma} [\delta_t(\tau) + a(\tau, \sigma)] + b(\sigma, x_{t+1}) \quad (6)$$

holds. The initialization is simply given by

$$\delta_0(\sigma) = 0, \quad \forall \sigma \in \Sigma$$

$$\begin{aligned} \delta_1(\sigma) &= \max_{\tau \in \Sigma} [\delta_0(\tau) + a(\tau, \sigma)] + b(\sigma, x_1) \\ &= a(\epsilon, \sigma) + b(\sigma, x_1). \end{aligned}$$

# The Inductive Step

The recursion step is given for  $2 \leq t \leq T$  by

$$\begin{aligned}
 \delta_t(\sigma) &= \max_{y_1, \dots, y_{t-1}} f(\xi, y_1, \dots, y_{t-1}, y_t = \sigma) \\
 &= \max_{y_1, \dots, y_{t-1}} \sum_{\tau, \bar{\tau} \in \mathcal{Y}} a(\tau, \bar{\tau}) \sum_{s=2}^{t-1} [[y_{s-1} = \tau \wedge y_s = \bar{\tau}]] \\
 &\quad + \sum_{\tau \in \Sigma} a(\tau, \sigma) [[y_{t-1} = \tau \wedge y_t = \sigma]] \\
 &\quad + \sum_{s=1}^{t-1} \sum_{\tau \in \Sigma} [[y_s = \tau]] b(\tau, x_s) + [[y_t = \sigma]] b(\sigma, x_t) \\
 &= \max_{\sigma^*} \max_{y_1, \dots, y_{t-2}} \sum_{\tau, \bar{\tau} \in \mathcal{Y}} a(\tau, \bar{\tau}) \sum_{s=2}^{t-2} [[y_{s-1} = \tau \wedge y_s = \bar{\tau}]] \\
 &\quad + \sum_{\tau \in \Sigma} a(\tau, \sigma^*) [[y_{t-2} = \tau \wedge y_{t-1} = \sigma^*]] \\
 &\quad + a(\sigma^*, \sigma) [[y^{t-1} = \sigma^* \wedge y^t = \sigma]] \\
 &\quad + \sum_{s=1}^{t-2} \sum_{\tau \in \Sigma} [[y_s = \tau]] b(\tau, x_s) + b(\sigma^*, x_{t-1}) + b(\sigma, x_t)
 \end{aligned}$$

# The Inductive Step Contd.

$$\begin{aligned} &= \max_{\sigma^*} \left[ \max_{y_1, \dots, y_{t-2}} f(\xi, y_1, \dots, y_{t-2}, y_{t-1} = \sigma^*) + a(\sigma^*, \sigma) \right] + b(\sigma, x_t) \\ &= \max_{\sigma^*} [\delta_{t-1}(\sigma^*) + a(\sigma^*, \sigma)] + b(\sigma, x_t). \end{aligned}$$

Thus, the top scoring sequence has the score

$$\max f(\xi, ) = \max_{\sigma \in \Sigma} \delta_T(\sigma).$$

We only sketch the extension to the argument of the maximum since it is analogous to the regular Viterbi algorithm. We introduce path variables  $\varphi_t(\sigma)$  that are initialized by  $\varphi_1(\sigma) = \epsilon$  for all  $\sigma \in \Sigma$ .

# Computing the Argmax

The sequence  $\varphi_t(\sigma)$  is then defined recursively for  $2 \leq t \leq T$  by

$$\varphi_t(\sigma) =_{\sigma^* \in \Sigma} [\delta_{t-1}(\sigma^*) + a(\sigma^*, \sigma)] .$$

Once the  $\delta_t(\sigma)$  of Theorem 2 are fixed, the optimal label sequence can be found by backtracking

$$y_T^* =_{\sigma \in \Sigma} \delta_T(\sigma)$$

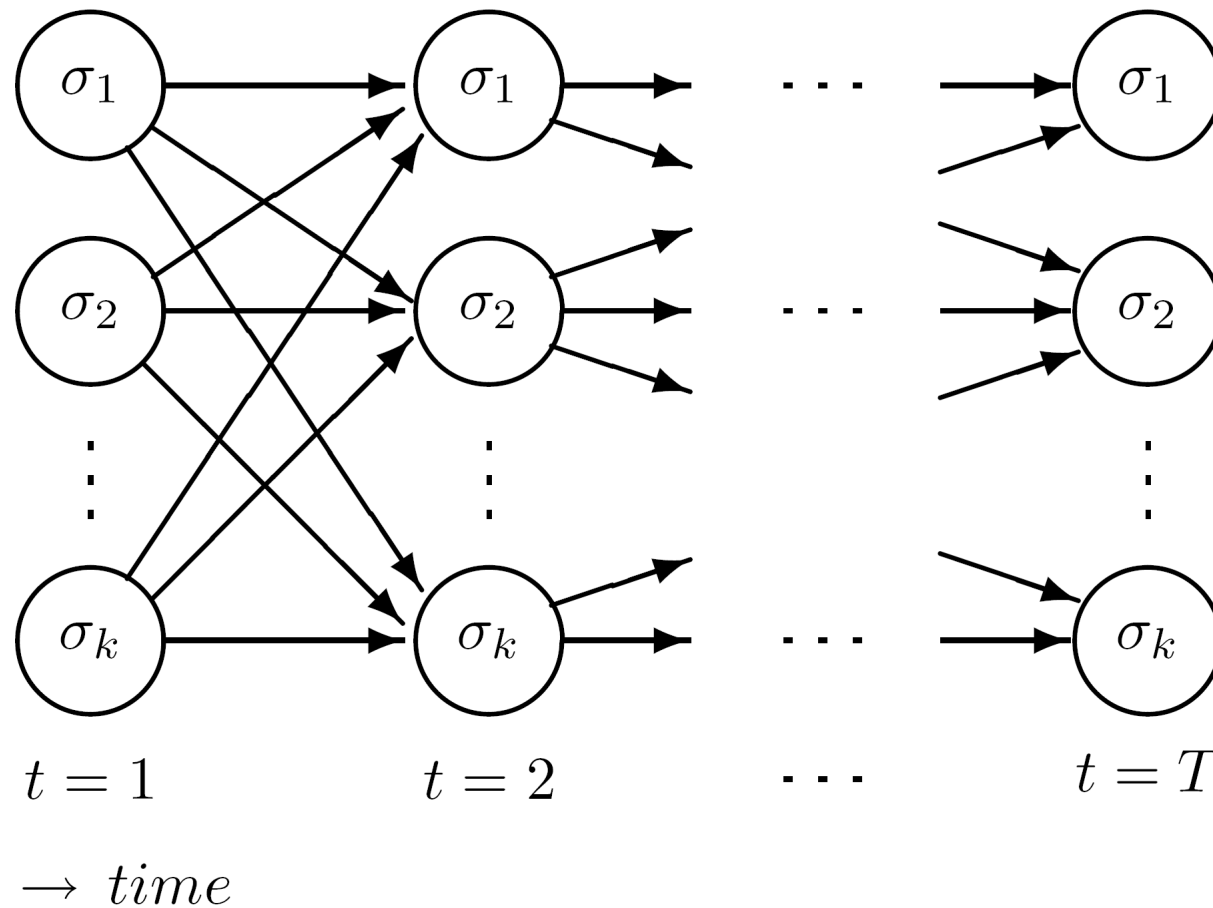
$$y_t^* = \varphi_{t+1}(y_{t+1}^*) \quad \text{for } t = T - 1, \dots, 1.$$



# Conclusion

Given the transition matrix  $[\mathbf{A}]_{\sigma,\tau} = a(\sigma, \tau)$  and the observation matrix  $[\mathbf{B}_\xi]_{\sigma,t} = b(\sigma, x_t)$  for input  $\xi$ , the computation of  $\delta$  and  $\varphi$  for a fixed  $t$  and  $\sigma \in \Sigma$  involves visiting  $|\Sigma|$  predecessors; thus, for a sequence of length  $T$  the time needed is in  $\mathcal{O}(T|\Sigma|^2)$ . This concludes the proof. □

# Visualization



**Figure:** Visualization of a trellis over the alphabet  $\Sigma = \{\sigma_1, \dots, \sigma_k\}$ .

# Summary

- Equivalence: Dual perceptron  $f(\xi, )$  and Viterbi algorithm
  - similar proof for primal perceptron
- pos:
  - easy to implement
  - efficient training process
- neg:
  - depends on ordering
  - no confidences
  - only 0/1 loss

# From Perceptrons to SVMs

- Add confidence to decision
- Incorporate arbitrary (structured) loss functions
- Impact of ordering resolved by quadratic programming

# Confidence Term

- Perceptron:

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, A, A]) \rangle$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, A, N]) \rangle$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, N, A]) \rangle$$

$$\vdots > \vdots$$

- Now, add a confidence  $\bar{\gamma}$ :

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle - \langle \wedge, \Phi(\xi, [A, A, A, A]) \rangle \geq \bar{\gamma}$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle - \langle \wedge, \Phi(\xi, [A, A, A, N]) \rangle \geq \bar{\gamma}$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle - \langle \wedge, \Phi(\xi, [A, A, N, A]) \rangle \geq \bar{\gamma}$$

$$\vdots \qquad \vdots$$

# Optimization Problem

$$\begin{aligned} \max_{\bar{\gamma}, \wedge} \quad & \frac{\bar{\gamma}}{\|\wedge\|} \\ \text{s.t.} \quad & \forall_{i=1}^n, \forall_{\neq i} : \langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \cdot) \rangle \geq \bar{\gamma} \end{aligned}$$

- We call
  - $\wedge$  the weight vector
  - $\bar{\gamma}$  the functional margin
  - $\gamma = \frac{\bar{\gamma}}{\|\wedge\|}$  the geometrical margin
- Problem:  $\bar{\gamma}$  and  $\wedge$  interdepend!
  - Remedy: fix one, solve for the other
  - Common approach:  $\bar{\gamma} = 1$ .

# Structural Hard-margin SVM

$$\begin{array}{ll} \min_{\wedge} & \frac{1}{2} || \wedge ||^2 \\ \text{s.t.} & \forall_{i=1}^n, \forall_{\neq i} : \langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle \geq 1 \end{array}$$

- Converges only when data is linear separable
- Remedy: allow for pointwise relaxations of the margin constraint
  - introduce slack variables  $\xi_i$  for input examples

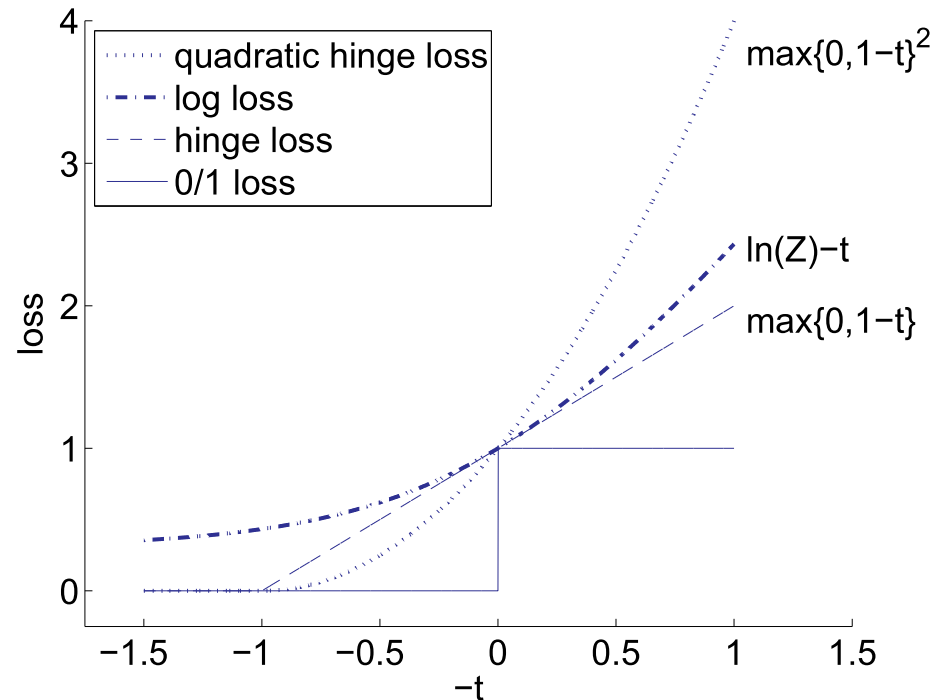
# Structural Soft-margin SVM

$$\begin{aligned}
 \min_{\wedge} \quad & \frac{1}{2} || \wedge ||^2 + \sum_{i=1}^n \xi_i \\
 \text{s.t.} \quad & \forall_{i=1}^n, \forall_{\neq i} : \langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle \geq 1 - \xi_i \\
 & \forall_{i=1}^n : \xi_i \geq 0
 \end{aligned}$$

- Sum of slacks upper bounds 0/1 loss
- Now: maximize margin between true  $i$  and best runner-up $^-$
- Alternative formulation:
  - slack  $\xi_i^-$  are bound to constraint  $\langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle$
  - computationally demanding



# Hinge-loss



- SVM implicitly implements a hinge loss (solve for slacks)
- Hinge loss can be rescaled to incorporate arbitrary loss functions
  - Let  $\Delta(i, \hat{y})$  denote a structural loss.
  - $\Delta : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}_0^+$ .
  - $\Delta(i, i) = 0$

# Exemplary Loss Functions

- 0/1 loss:  $\Delta(, \bar{\cdot}) = [[==]]$
- Hamming loss for sequences

$$\Delta(, \bar{\cdot}) = T - \sum_{t=1}^T [[y_t == \bar{y}_t]]$$

- Property: decomposes across the cliques!

# Margin-rescaling

- Taskar et al. (2004)
- Rescale the (functional) margin by actual loss

$$\begin{aligned} \min_{\wedge} \quad & \frac{1}{2} \|\wedge\|^2 + \sum_{i=1}^n \xi_i \\ \text{s.t.} \quad & \forall_{i=1}^n, \forall_{\neq i} : \langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_{i,\neq}) \rangle \geq \Delta(i,\neq) - \xi_i \\ & \forall_{i=1}^n : \xi_i \geq 0 \end{aligned}$$

- Implicit hinge loss upper bounds  $\Delta$
- Most strongly violated constraint:

$$-\left(\Delta(i,\neq) - (\langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_{i,\neq}) \rangle)\right)$$

# Slack-rescaling

- Tsochantaridis et al. (2005)
- Rescale slack variables by actual loss

$$\begin{aligned}
 \min_{\wedge} \quad & \frac{1}{2} \|\wedge\|^2 + \sum_{i=1}^n \xi_i \\
 \text{s.t.} \quad & \forall_{i=1}^n, \forall_{\neq i} : \langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle \geq 1 - \frac{\xi_i}{\Delta(i, \neg)} \\
 & \forall_{i=1}^n : \xi_i \geq 0
 \end{aligned}$$

- Implicit hinge loss upper bounds  $\Delta$
- Most strongly violated constraint:

$$- \left( 1 - \Delta(i, \neg) \times (\langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle) \right)$$

# Implications

- Loss  $\Delta$  decomposes across the cliques of the graph
  - Margin-rescaling is easily integrated into inference
  - Slack-rescaling difficult
- Loss not decomposable
  - Both difficult!
- In practice, slack-rescaling often better than margin-rescaling
  - rarely applicable (needs good approximation or enumerable sets)

# Recall:

- relation between
  - $P(|\xi)$
  - model  $f(\xi, )$
  - log-Viterbi algorithm
- intuition
  - $f(\xi, ) =$  how good does fits to  $\xi$
  - log-Viterbi: find top-scoring

# Towards Structured Support Vector Machines

- Add confidence to decision
- Incorporate arbitrary (structured) loss functions
- Impact of ordering resolved by quadratic programming

# Confidence Term

- Perceptron:

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, A, A]) \rangle$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, A, N]) \rangle$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle > \langle \wedge, \Phi(\xi, [A, A, N, A]) \rangle$$

$$\vdots > \vdots$$

- Now, add a confidence  $\bar{\gamma}$ :

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle - \langle \wedge, \Phi(\xi, [A, A, A, A]) \rangle \geq \bar{\gamma}$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle - \langle \wedge, \Phi(\xi, [A, A, A, N]) \rangle \geq \bar{\gamma}$$

$$\langle \wedge, \Phi(\xi, [N, V, A, N]) \rangle - \langle \wedge, \Phi(\xi, [A, A, N, A]) \rangle \geq \bar{\gamma}$$

$$\vdots \qquad \vdots$$



# Optimization Problem

$$\begin{aligned} \max_{\bar{\gamma}, \wedge} \quad & \frac{\bar{\gamma}}{\|\wedge\|} \\ \text{s.t.} \quad & \forall_{i=1}^n, \forall_{\neq i} : \langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \cdot) \rangle \geq \bar{\gamma} \end{aligned}$$

- We call
  - $\wedge$  the weight vector
  - $\bar{\gamma}$  the functional margin
  - $\gamma = \frac{\bar{\gamma}}{\|\wedge\|}$  the geometrical margin
- Problem:  $\bar{\gamma}$  and  $\wedge$  interdepend!
  - Remedy: fix one, solve for the other
  - Common approach:  $\bar{\gamma} = 1$ .

# Structural Hard-margin SVM

$$\begin{aligned} \min_{\wedge} \quad & \frac{1}{2} || \wedge ||^2 \\ \text{s.t.} \quad & \forall_{i=1}^n, \forall_{\neq i} : \langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle \geq 1 \end{aligned}$$

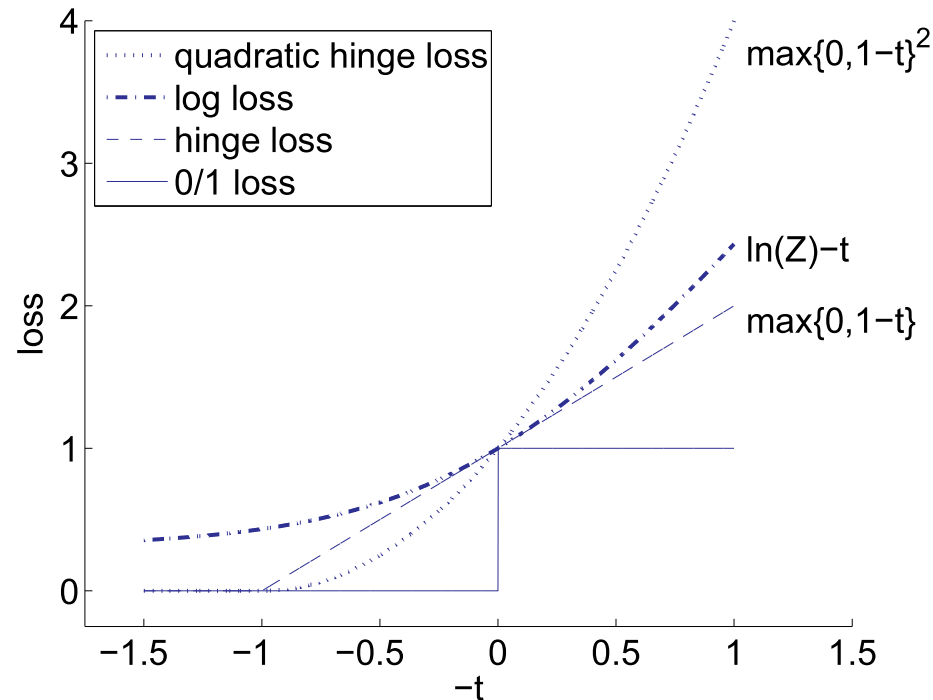
- Converges only when data is linear separable
- Remedy: allow for pointwise relaxations of the margin constraint
  - introduce slack variables  $\xi_i$  for input examples

# Structural Soft-margin SVM

$$\begin{aligned}
 \min_{\wedge} \quad & \frac{1}{2} || \wedge ||^2 + C \sum_{i=1}^n \xi_i \\
 \text{s.t.} \quad & \forall_{i=1}^n, \forall_{\neq i} : \langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle \geq 1 - \xi_i \\
 & \forall_{i=1}^n : \xi_i \geq 0
 \end{aligned}$$

- Maximize margin between true  $i$  and best runner-up<sup>-</sup>
  - Sum of slacks upper bounds 0/1 loss
  - Trade-off parameter  $C > 0$
- Alternative formulation:
  - slack  $\xi_i$  are bound to constraint  $\langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle$
  - computationally demanding

# Hinge-loss



- SVM implicitly implements a hinge loss (solve for slacks)
- Hinge loss can be rescaled to incorporate arbitrary loss functions
  - Let  $\Delta(i, \hat{y})$  denote a structural loss.
  - $\Delta : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}_0^+$ .
  - $\Delta(i, i) = 0$

# Exemplary Loss Functions

- 0/1 loss:  $\Delta(, \bar{\cdot}) = [[==]]$
- Hamming loss for sequences

$$\Delta(, \bar{\cdot}) = T - \sum_{t=1}^T [[y_t == \bar{y}_t]]$$

- Property: decomposes across the cliques!

# Risk Minimization

- We want to minimize the theoretical risk (the generalization error)

$$R(f) = \int_{\mathcal{X} \times \mathcal{Y}} \Delta(\cdot, f(\xi, \cdot)) dP(\xi, \cdot)$$

- In general, we don't know  $P(\xi, \cdot)$ 
  - Remedy: Use training sample instead!
- Minimize the empirical risk

$$\hat{R}(f) = \sum_{i=1}^n \Delta(\cdot, f(\xi_i, \cdot))$$

# Idea

- SVMs minimize the (regularized) empirical risk:

$$\hat{R}(f) = \sum_{i=1}^n \Delta(\xi_i, f(\xi_i, \cdot))$$

- Sum of slacks  $\sum_i \xi_i$  upper bounds empirical risk
- Slack variable  $\xi_i$  denotes the error for input  $\xi_i$ 
  - Now: Find maximal error wrt  $\Delta$

# Margin-rescaling

- Taskar et al. (2004)
- Rescale the (functional) margin by actual loss

$$\begin{aligned} \min_{\wedge} \quad & \frac{1}{2} \|\wedge\|^2 + C \sum_{i=1}^n \xi_i \\ \text{s.t.} \quad & \forall_{i=1}^n, \forall_{\neq i} : \langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle \geq \Delta(i, \neg) - \xi_i \\ & \forall_{i=1}^n : \xi_i \geq 0 \end{aligned}$$

- Implicit hinge loss upper bounds  $\Delta$
- Most strongly violated constraint:

$$\neg_{\neq i} \left( \underbrace{\Delta(i, \neg) - (\langle \wedge, \Phi(\xi_{i,i}) - \Phi(\xi_i, \neg) \rangle)}_{\xi_i} \right)$$



# Slack-rescaling

- Tsochantaridis et al. (2005)
- Rescale slack variables by actual loss

$$\begin{aligned}
 \min_{\wedge} \quad & \frac{1}{2} \|\wedge\|^2 + C \sum_{i=1}^n \xi_i \\
 \text{s.t.} \quad & \forall_{i=1}^n, \forall_{\neq i} : \langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle \geq 1 - \frac{\xi_i}{\Delta(i, \neg)} \\
 & \forall_{i=1}^n : \xi_i \geq 0
 \end{aligned}$$

- Implicit hinge loss upper bounds  $\Delta$
- Most strongly violated constraint:

$$\neg_{\neq i} \left( \underbrace{1 - \Delta(i, \neg) \times (\langle \wedge, \Phi(\xi_{i,i}) \rangle - \langle \wedge, \Phi(\xi_i, \neg) \rangle)}_{\xi_i} \right)$$

# Implications

- Loss  $\Delta$  decomposes across the cliques of the graph
  - Margin-rescaling is easily integrated into inference
  - Slack-rescaling difficult
- Loss not decomposable
  - Both difficult!
- In practice, slack-rescaling often better than margin-rescaling
  - rarely applicable (needs good approximation or enumerable sets)

# Example

- Margin rescaling for sequences / Viterbi algorithm
- Remainder: 0/1 loss for simplicity

# Towards Dual SVMs

- Integrate constraints into objective
  - Apply Lagrange's Theorem
  - Lagrange multipliers  $\alpha_i(\cdot)$  and  $\mu_i$

- Build Lagrangian:

$$\begin{aligned}
 L = & \frac{1}{2} \| \wedge \|^2 + C \sum_{i=1}^n \xi_i \\
 & - \sum_{i=1}^n \sum_{\vec{\neq}_i} \alpha_i(\cdot) \langle \wedge, \Phi(\xi_i, \vec{\neq}_i) \rangle - \langle \wedge, \Phi(\xi_i, \vec{\neq}_i) \rangle - 1 + \xi_i \\
 & - \sum_{i=1}^n \beta_i \xi_i
 \end{aligned}$$

- Minimum of Lagrangian is a saddle-point
  - max wrt  $\alpha, \mu$ , min wrt  $\wedge, \xi$

# Partial Derivatives: $\xi_i$

- Compute partial derivatives wrt  $\xi$ :

$$\frac{\partial L}{\partial \xi_i} = C - \sum_{\neq i} \alpha_i() - \beta_i \stackrel{!}{=} 0 \quad (7)$$

- Using the non-negativity of  $\alpha$  and  $\beta$  yields

$$\forall_{i=1}^n : \quad 0 \leq \sum_{\neq i} \alpha_i() \leq C$$

# Partial Derivatives: $\Lambda$

- Compute partial derivatives wrt  $\Lambda$ :

$$\frac{\partial L}{\partial \Lambda} = \mathbf{w} - \sum_{i=1}^n \sum_{\neg \neq_i} \alpha_i(\neg) \left( \Phi(\xi_{i,i}) - \Phi(\xi_i, \neg) \right) \stackrel{!}{=} 0.$$

- We obtain:

$$\Lambda = \sum_{i=1}^n \sum_{\neg \neq_i} \alpha_i(\neg) \left( \Phi(\xi_{i,i}) - \Phi(\xi_i, \neg) \right)$$

- Recall: dual perceptron!
  - Definition of  $\Lambda$  is equivalent
  - $\alpha$ 's act like counters